

I320D – Topics in Human Centered Data Science
Text Mining and NLP Essentials


Week 12: Language Models and Embeddings (...) , NLP Applications

Dr. Abhijit Mishra

Ongoing Assignments / Project

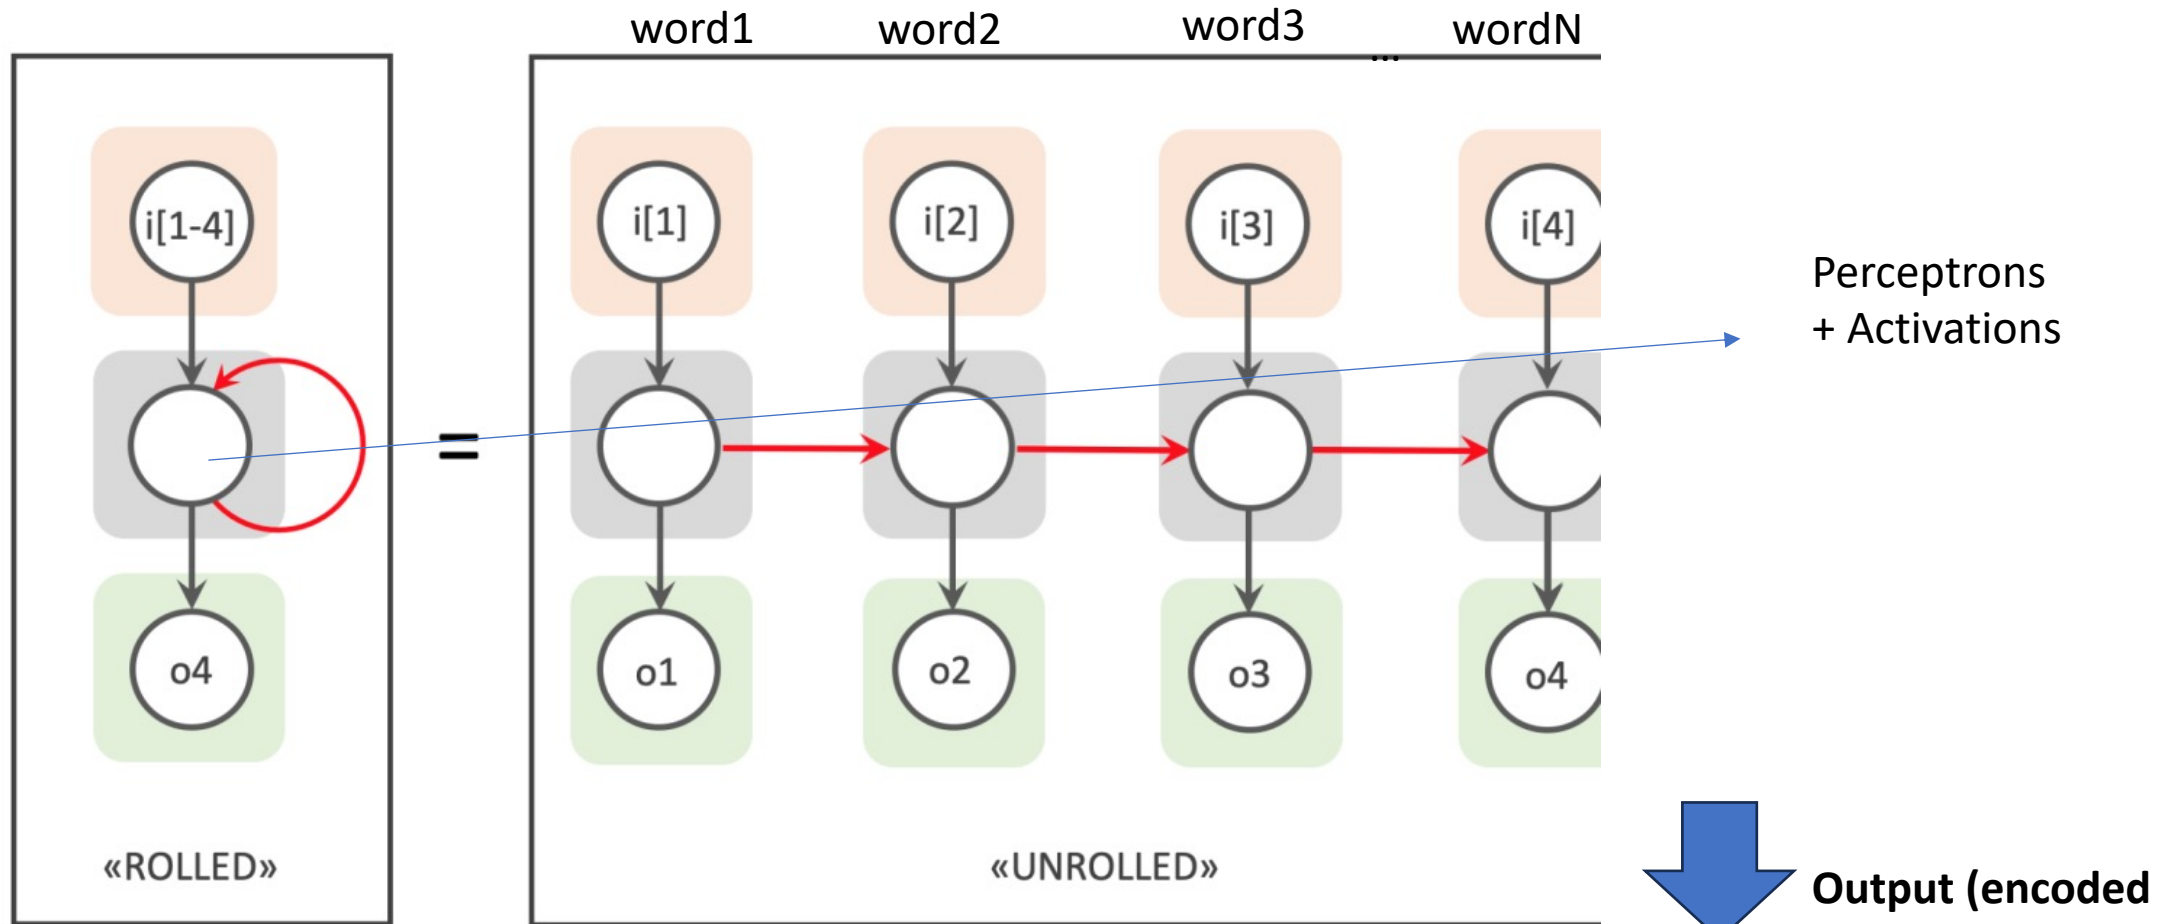
- **Course Project:**
 - Feedback Shared, leverage office hours on Monday / Wednesday
 - 5 mins “Work progress” presentation on **Apr 15**
 - Final Presentation: **Apr 29**
 - Final Report Due : **May 6**
- **Assignment 5: Text Classification (Due April 7)**

Wee 11 Recap: Feed Forward Netwrks



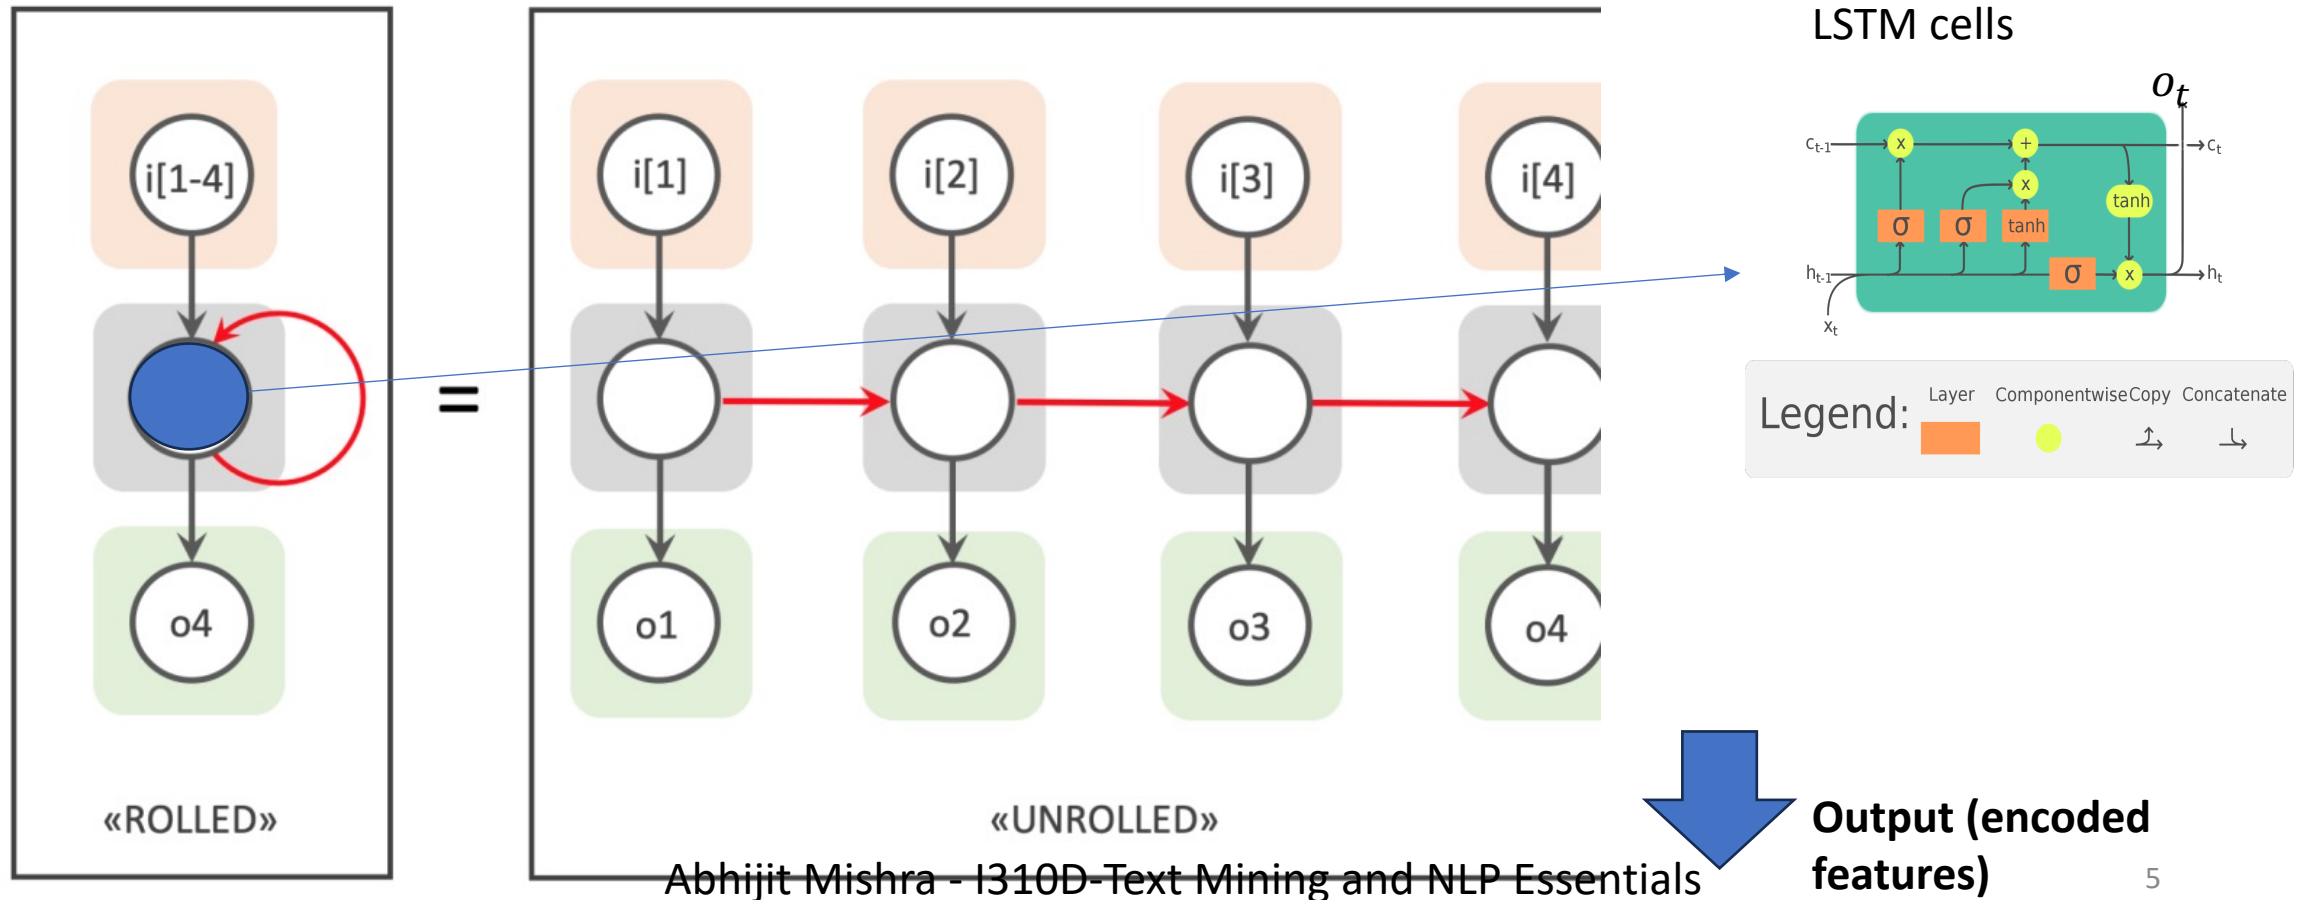
Week 11: Recap

Recurrent Neural Nets (RNNs)



Week 11: Recap

Long Short Term Memories (Schmidhuber et al, 1997)



In Python

```
# Define the LSTM model
```

```
model = Sequential()
```

```
model.add(Embedding(max_features, 128, input_length=maxlen))
```

```
model.add(SpatialDropout1D(0.2))
```

```
model.add(LSTM(128, dropout=0.2, recurrent_dropout=0.2))
```

```
model.add(Dense(1, activation='sigmoid'))
```

```
# Compile the model
```

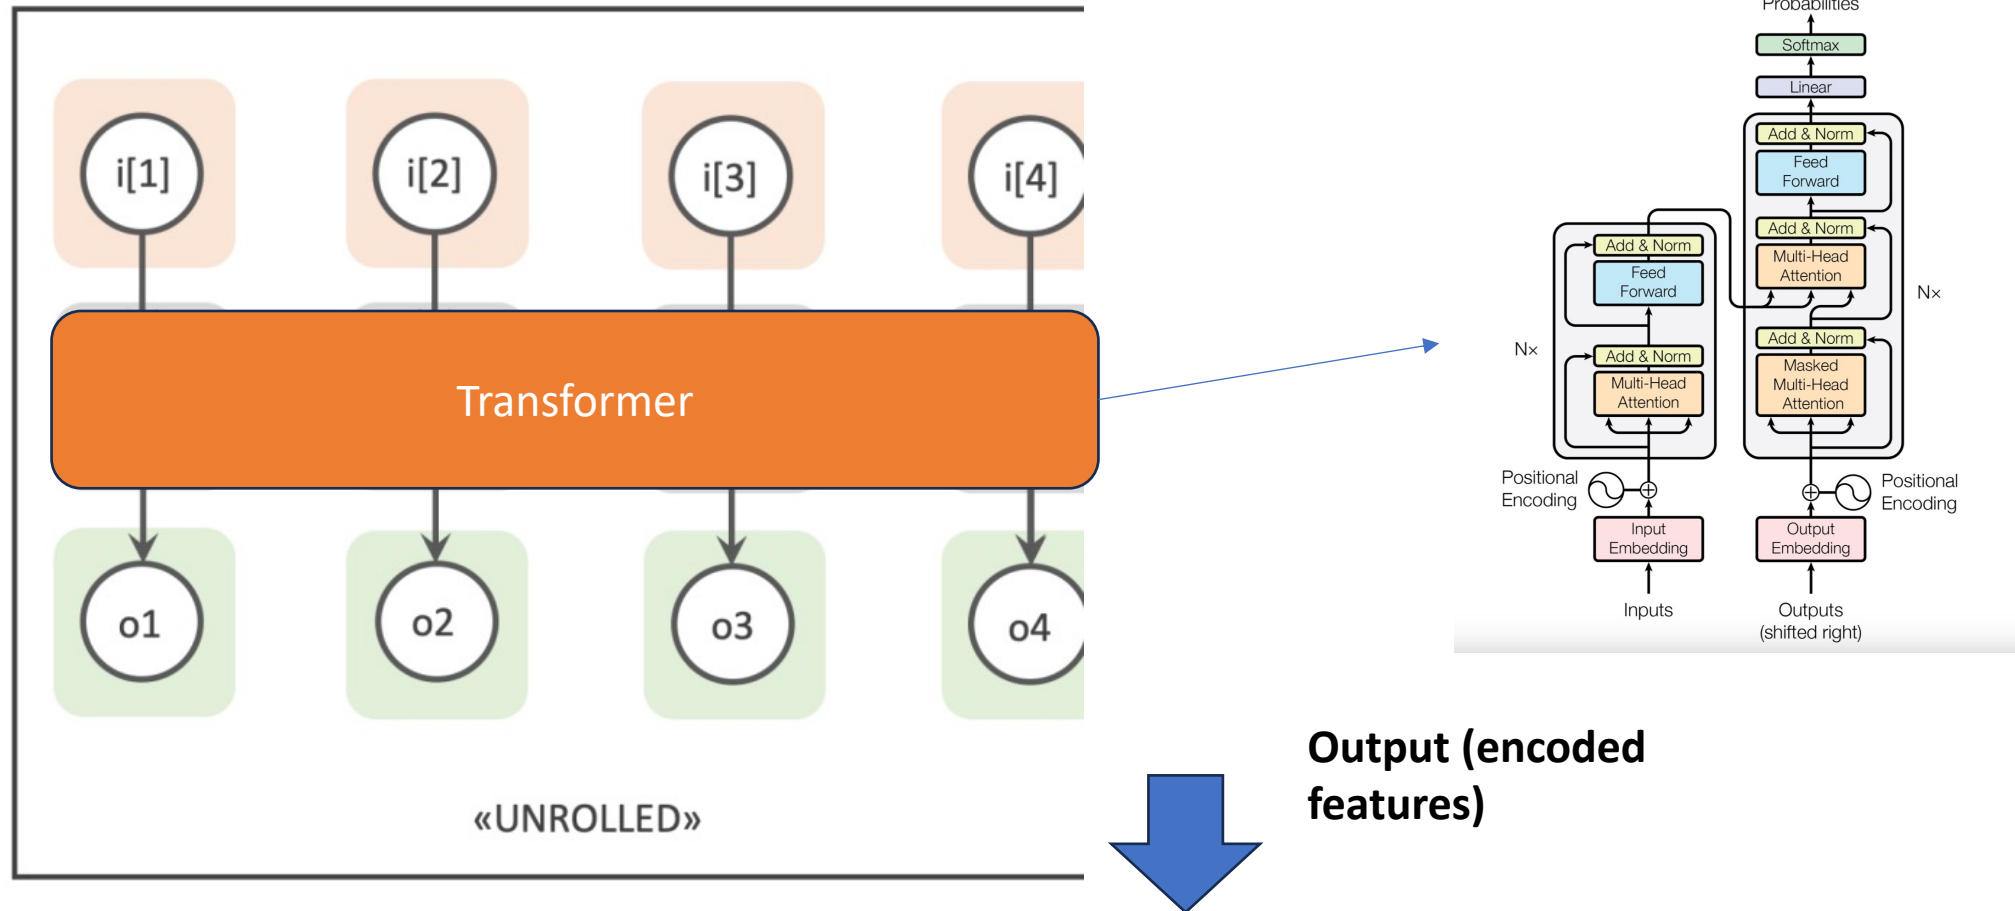
```
model.compile(loss='binary_crossentropy', optimizer='adam',  
metrics=['accuracy'])
```

```
# Train the model
```

```
model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs,  
validation_data=(x_test, y_test))
```

Week 11: Recap

Transformers (Vaswani et al, 2018)



Recap: Language Models (LMs)

- Language models are statistical or deep learning models that learn to predict the probability of a sequence of words in a sentence or text
- For a sequence of words $W = (w_1, w_2, w_3, \dots, w_n)$
- A language model can be expressed as

$$f(X, \theta) \rightarrow \frac{P(W|\theta) = P(w_1|\theta) \cdot P(w_2|w_1, \theta) \cdot P(w_3|w_1, w_2, \theta) \cdot \dots \cdot P(w_n|w_1, w_2, \dots, w_{n-1}, \theta)}{P(w_n|w_1, w_2, \dots, w_{n-1}, \theta)}$$

- Here theta => model parameters

LMs are Generative Models

- Language Models are Generative in Nature

Generative Modeling of Text

- Tasks where:

- Input is a sequence $X = \{x_1, x_2 \dots, x_N\}$ **or** $\mathbf{X} \in \mathbb{R}^N$
- Output is a sequence $Y = \{y_1, y_2 \dots, y_M\}$ **or** $\mathbf{Y} \in \mathbb{R}^M$

- Example:

- Text summarization
- Machine Translation
- Chat generation

Sequence Generation – Text Summarization Example

SOURCE: *Roger Federer wins a record eighth men's singles title at Wimbledon on Sunday.*

TARGET:

Roger Federer won the Wimbledon



REPRERSENTATION

Sequence Generation – English-Spanish Translation Example

SOURCE: *Roger Federer wins a record eighth men's singles title at Wimbledon on Sunday.*

TARGET:
Roger Federer gana un octavo título individual masculino en Wimbledon el domingo.




REPRERSENTATION

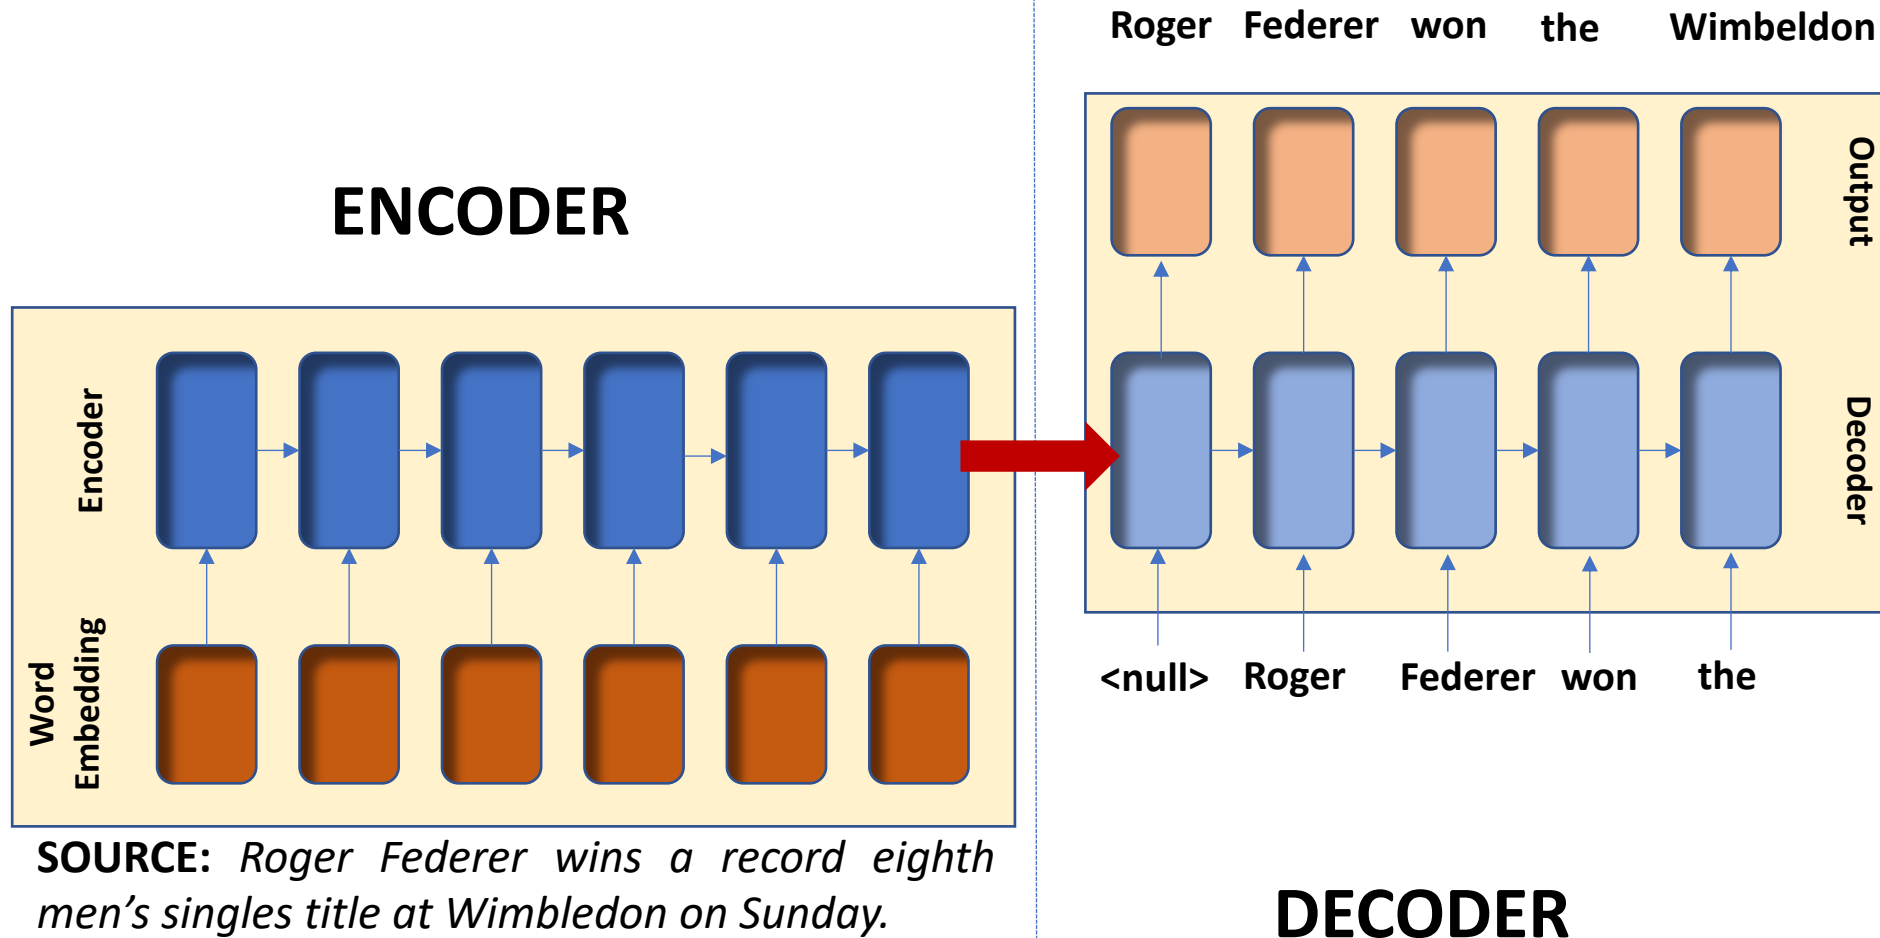
Sequence Generation – Language Modeling Example

SOURCE: *Roger Federer wins a record eighth*

TARGET:
men's singles title at Wimbledon on Sunday.



Zooming into Encoder-Decoder Models




We have Various choices for the BLUE blocks (RNNs, LSTMs, Transformers)

Embeddings: Representation Learning

- **Feature Engineering from input words / input sentences?**
 - What do we intend to do?
- Extract meaningful representations:
 - In computer understandable numerical forms
 - Should capture relationships between words
 - Synonymy (e.g., “specimen”, “sample”)
 - Antonymy (e.g., ”man”, “woman”)
 - Conceptual similarity (e.g., “Wednesday”, “Monday”) (“USA”, ”Canada”)

Issues with One-hot vector

Word	1-hot vector
Queen	[1, 0, 0]
King	[0, 1, 0]
Cat	[0, 0, 1]

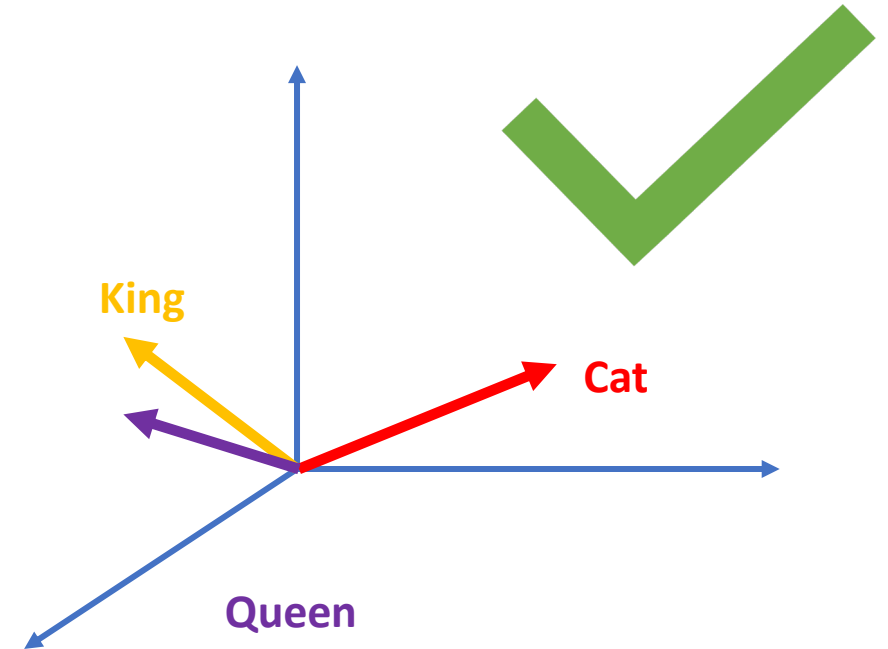


$$D_{\cosine}(\textit{Cat}, \textit{King}) = D_{\cosine}(\textit{Cat}, \textit{Queen}) = D_{\cosine}(\textit{King}, \textit{Queen}) = 1$$

$$D_{Euclid}(\textit{Cat}, \textit{King}) = D_{Euclid}(\textit{Cat}, \textit{Queen}) = D_{Euclid}(\textit{King}, \textit{Queen}) = \sqrt{2}$$

Instead, we need

Word	1-hot vector
Queen	[1.5, -1.3, -0.9]
King	[2.1, -0.7, 0.2]
Cat	[0.3, 1.9, -0.4]



$$D_{\cosine}(\textit{Cat}, \textit{King}) = 1.17, D_{\cosine}(\textit{Cat}, \textit{Queen}) = 1.38,$$

$$D_{\cosine}(\textit{King}, \textit{Queen}) = 0.19$$

$$D_{Euclid}(\textit{Cat}, \textit{King}) = 3.21, D_{Euclid}(\textit{Cat}, \textit{Queen}) = 3.45$$

$$D_{Euclid}(\textit{King}, \textit{Queen}) = 1.38$$

What are Word Embeddings?

- Embeddings are matrices of shape $V \times E$.
- V represents the size of the vocabulary (or total number of valid words in a language).
- E represents the dimension of each vector for the word.
- Typically $E \ll V$
- In other words, we are projecting sparse **one-hot encodings** of words (of dimension V) to dense Embeddings of size E
- Size of E is typically in 100s (300, 600, 1000).

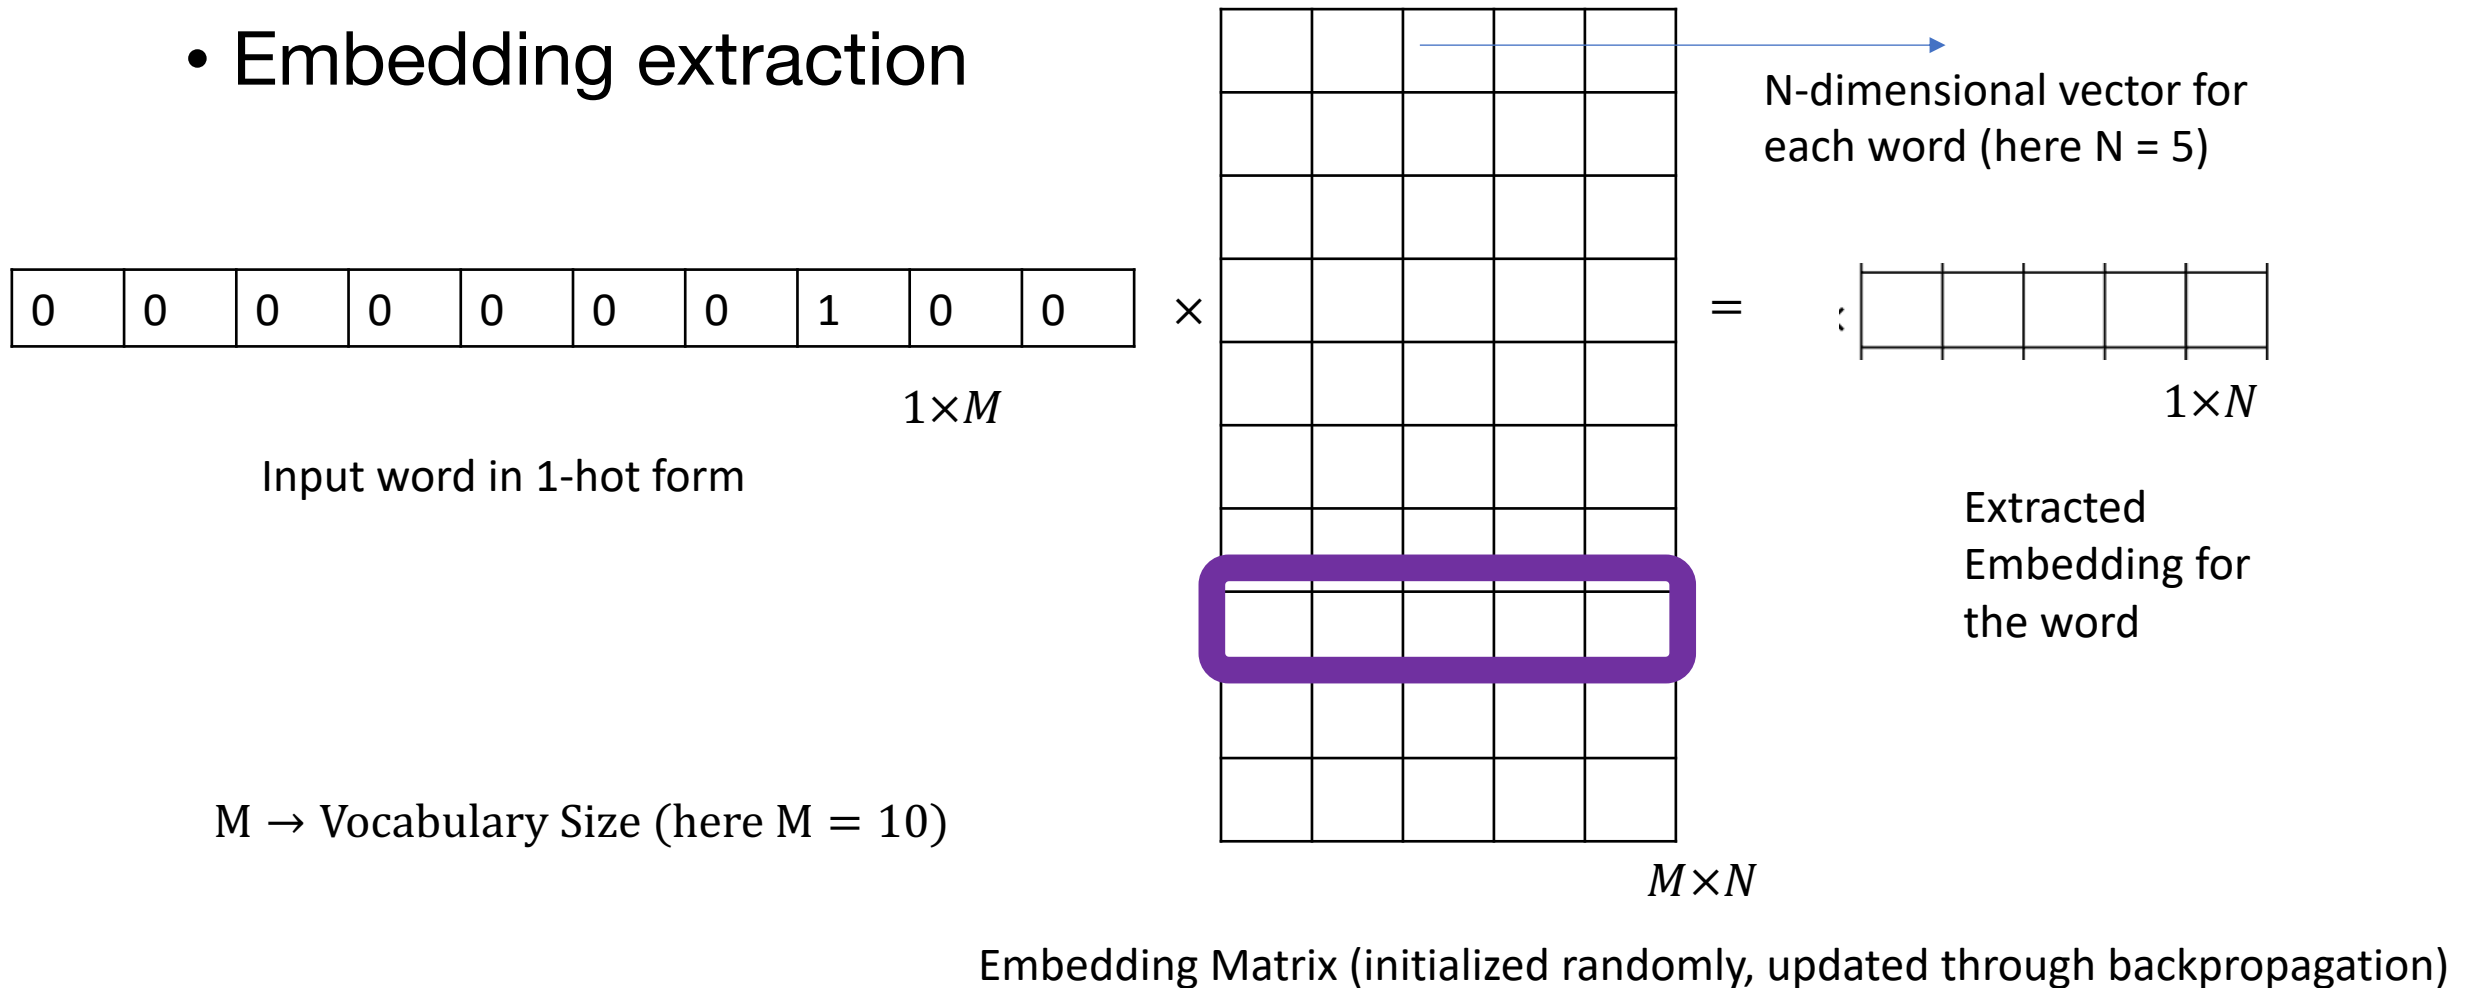
Digression: Matrix Multiplication

- Multiplying two matrices of shape $M \times N$ and $N \times P$ yields an $M \times P$ matrix
- Example:

$$\begin{aligned} [1 \ 2 \ 3] \times \begin{bmatrix} 2 & 3 \\ 4 & 5 \\ 6 & 7 \end{bmatrix} &= [1 \times 2 + 2 \times 4 + 3 \times 6 \quad 1 \times 3 + 2 \times 5 + 3 \times 7] \\ &= [2 + 8 + 18 \quad 3 + 10 + 21] = [28 \ 34] \end{aligned}$$


Digression: Projection through matrix multiplication

- Embedding extraction




Encoder Decoder Models with Embeddings

- Example Sentence: **“The cat sat on the mat”**




- Processing: “**The**”




Trainable embedding matrix

- Processing: “**cat**”




Trainable embedding matrix

- Processing: “**sat**”



Trainable embedding matrix

- Processing: “...”



Trainable embedding matrix

Embedding Layer

- A matrix of size (*Vocabulary size* \times *Embedding dimension*)
- Initialized with random values but updated using backpropagation
- From 1-hot to embeddings
 - 1 lookup operation

Word2Vec: Learning Embeddings with Feed Forward Encoder Decoder based LMs

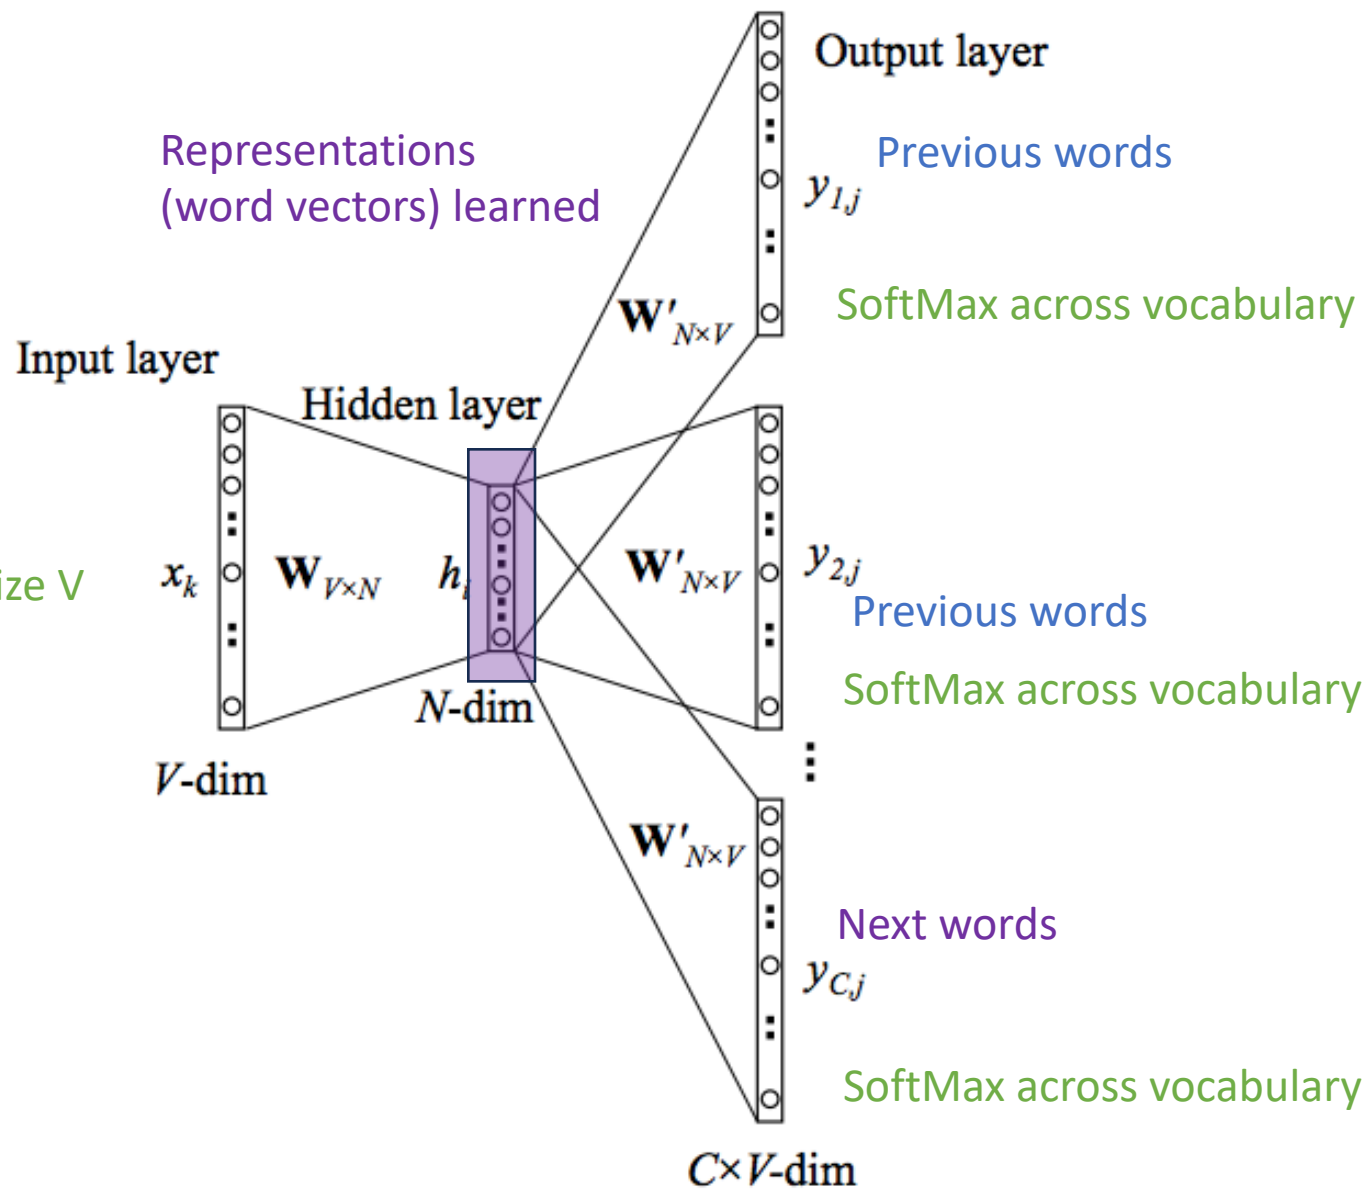
- **Skip Gram Objective** : Given a word can we predict the previous and the next words (or predict surrounding context given an input
- A feed forward network can be designed to perform this task
- **Dataset:**
 - Examples containing <input word, context> can be automatically created using large amount of corpus (e.g., Wikipedia, News Database etc)

Skip Gram Example

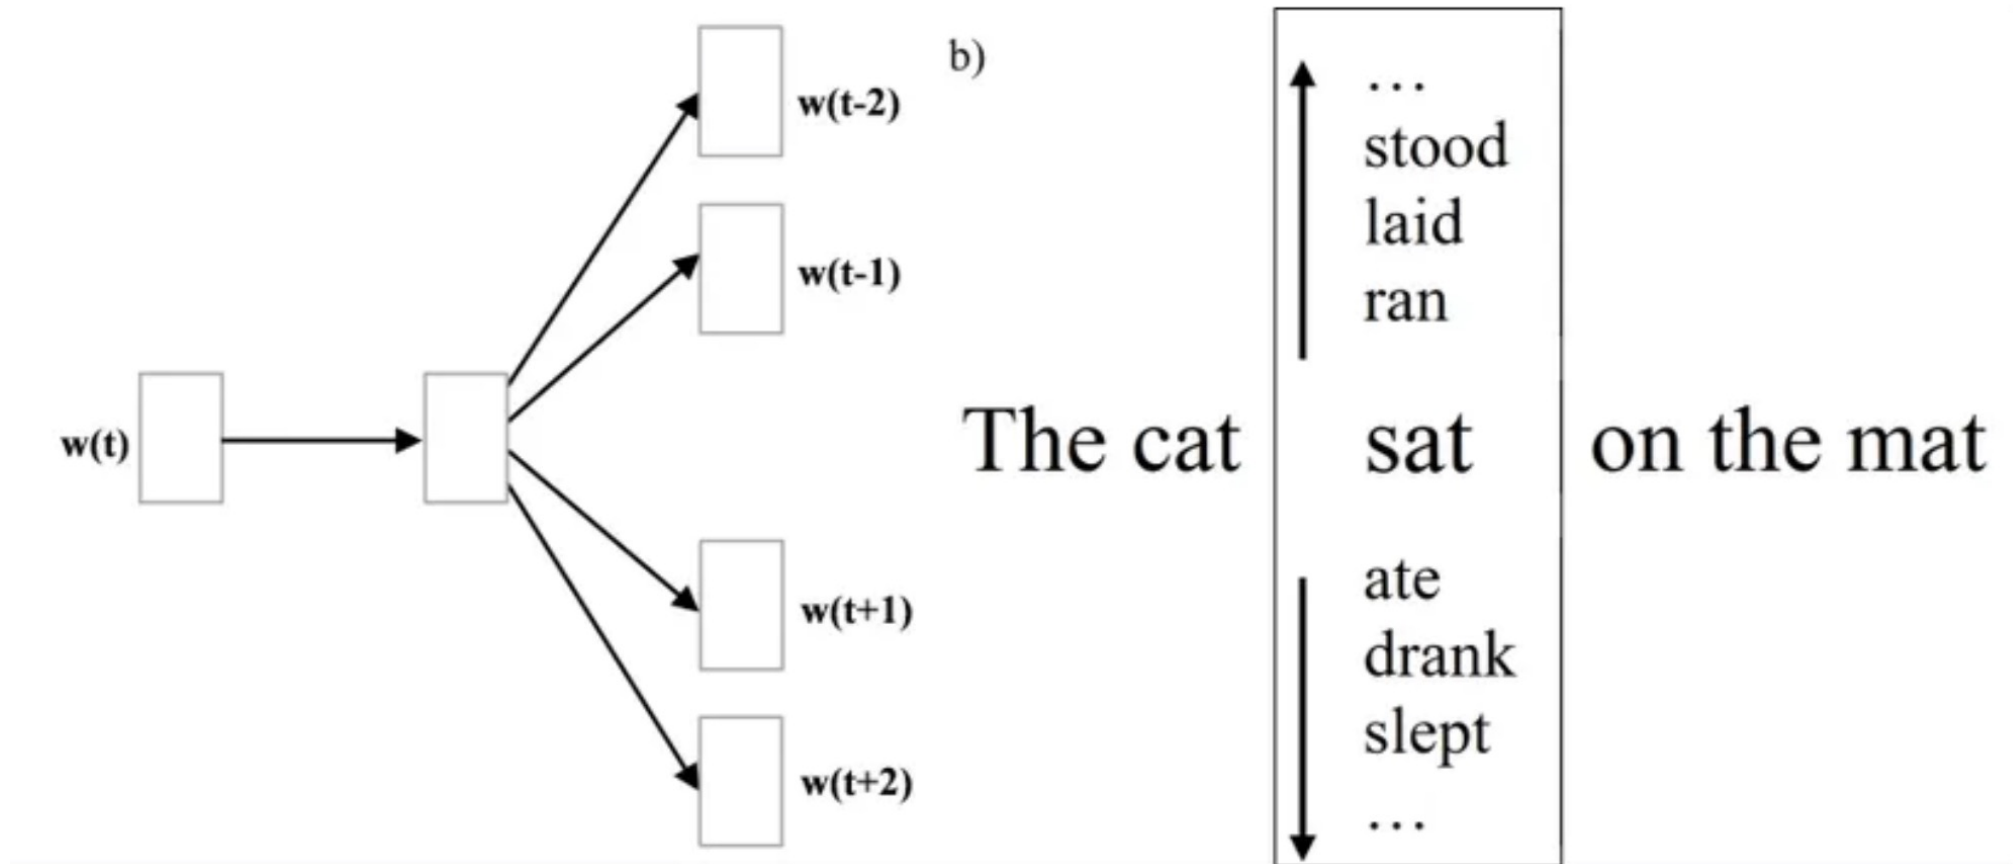
- Original sentence: “The cat sat on the mat”
- Preparing Training data:
 - **Input: “sat”**
 - **Output: “the cat on”**

SkipGram Model using Feed Forward Nets


Input word (1-hot vector) of size V



SkipGram Model using Feed Forward Nets



After Training: During Inference



Another Option: The CBOW Model

- What about we predict center word, **given context word, opposite to the skip-gram model?**
- Yes, this is called Continuous Bag Of Words model in the original Word2Vec paper.

Word2Vec: Pros and Cons

- **Pros:**

- Respects language and order to some extent
- Efficient Training Process: Simple FFDs
- Semantic Relationships Preservation

- **Cons:**

- Loss of local context
- Does not capture POS variations and word senses
 - Word “bank” will mostly be treated as NOUN
 - Word “bank” will always yield the same vector

GloVE

GloVe: Global Vectors for Word Representation

Jeffrey Pennington, Richard Socher, Christopher D. Manning

Computer Science Department, Stanford University, Stanford, CA 94305

`jpennin@stanford.edu, richard@socher.org, manning@stanford.edu`

- The GloVe algorithm extracts word vectors by optimizing a defined objective function that captures the statistical co-occurrence information between words

Glove-step-2: Form objective function

$$\text{minimize}_{w_1, w_2, \dots, w_V, b_1, b_2, \dots, b_V} (J)$$

$$J = \sum_{i,j=1}^V f(X_{ij}) \left(\overbrace{w_i^T \tilde{w}_j + b_i + \tilde{b}_j}^{\text{Parameters to be learned}} - \log X_{ij} \right)^2$$

Co-occurrence between word i and j

Parameters to be learned


Co-occurrence between word i and j

$f(X_{ij})$ is a weighting function that assigns more weight to less frequent co-occurrences to prevent extremely frequent words from dominating the training.

Glove-step-2: Optimize objective

$$\text{minimize}_{w_1, w_2, \dots, w_V, b_1, b_2, \dots, b_V} (J)$$

$$J = \sum_{i,j=1}^V f(X_{ij}) \left(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij} \right)^2$$




**These are the word
embeddings and
needs to be learned**

How?

$$\text{minimize}_{w_1, w_2, \dots, w_V, b_1, b_2, \dots, b_V} (J)$$

$$J = \sum_{i,j=1}^V f(X_{ij}) \left(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij} \right)^2$$




**Initialize randomly
and update through
gradient descent**

How?

$$\text{minimize}_{w_1, w_2, \dots, w_V, b_1, b_2, \dots, b_V} (J)$$

$$J = \sum_{i,j=1}^V f(X_{ij}) \left(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij} \right)^2$$



We can randomly initialize a d dimensional vector per word (e.g., d = 50, d = 200)

How to evaluate word embeddings?

- Also by word analogy tasks (Mikolov et al, 2013)
- Solving analogies of the form "**a is to b as c is to ?**" or "**a:b :: c:**" where you are given three words and you need to find the fourth word that completes the analogy
- Examples:
 1. "Man is to woman as king is to ____"
 2. "Spain is to Madrid as France is to ____"
 3. "Eat is to food as drink is to ____"

Sentence Vectors

Transformer Based LM: BERT Example

- BERT: **B**idirectional **E**ncoder **R**epresentation **T**ransformers
- Trained with two objectives :
 - Masked token prediction
 - Next sentence prediction

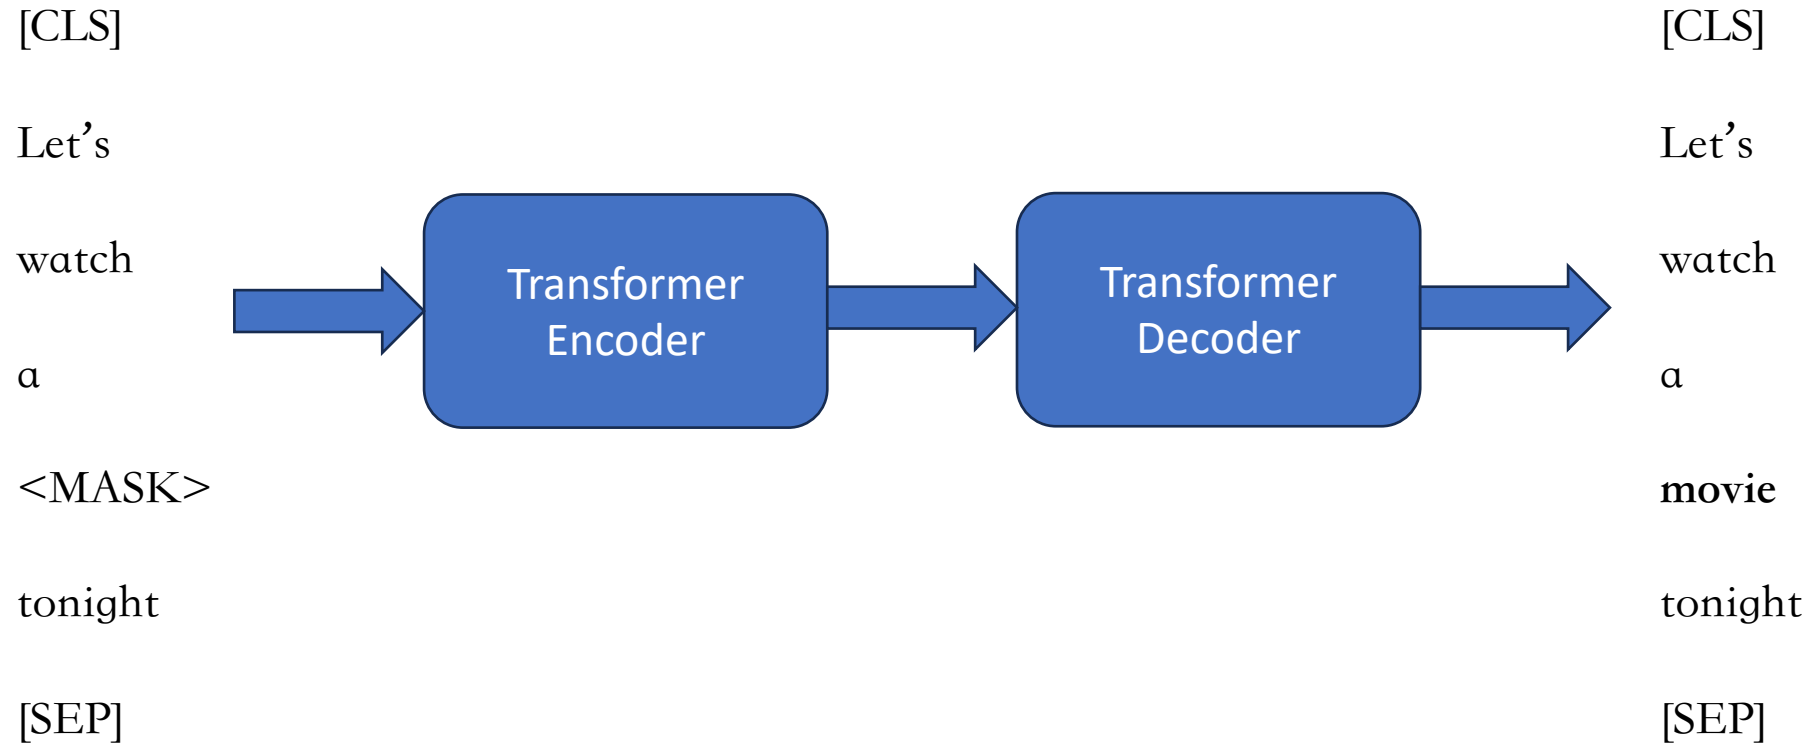
BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin Ming-Wei Chang Kenton Lee Kristina Toutanova

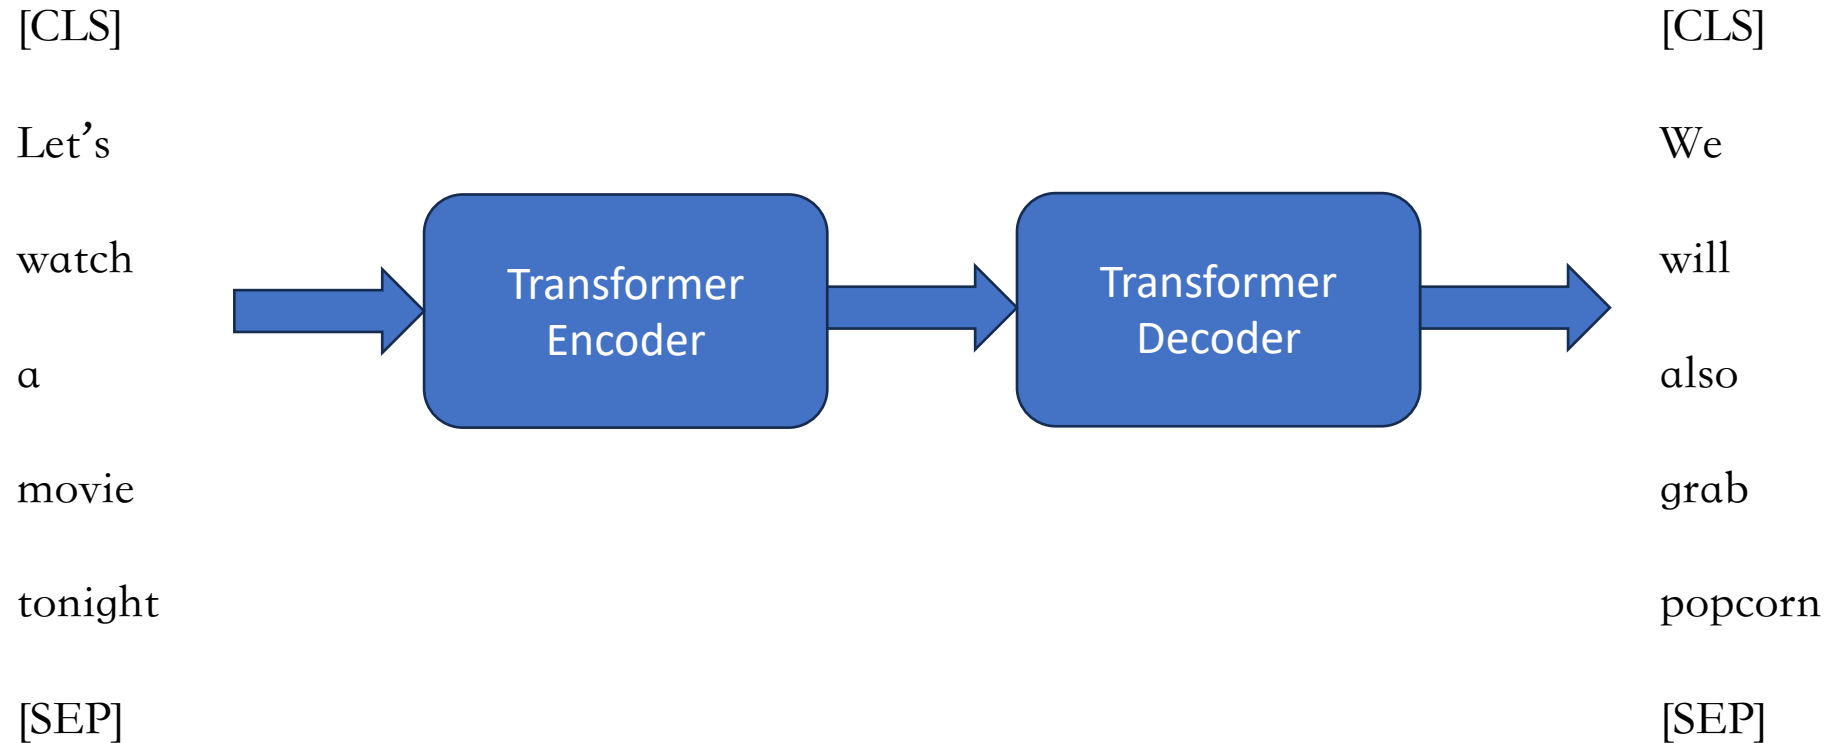
Google AI Language

`{jacobdevlin, mingweichang, kentonl, kristout}@google.com`

Training Task: Masked Token Prediction



Training Task: Next Sentence Prediction



Transformer Based LM: BERT Example

[CLS]

Let's


watch

a

movie

tonight

[SEP]



Sentence Embeddings Extraction /
Sentence Feature Extraction

**We only use a portion of a network that helps extract features
(also known as encoder)**

Sentence Vectors: Pros and Cons

- **Pros:**

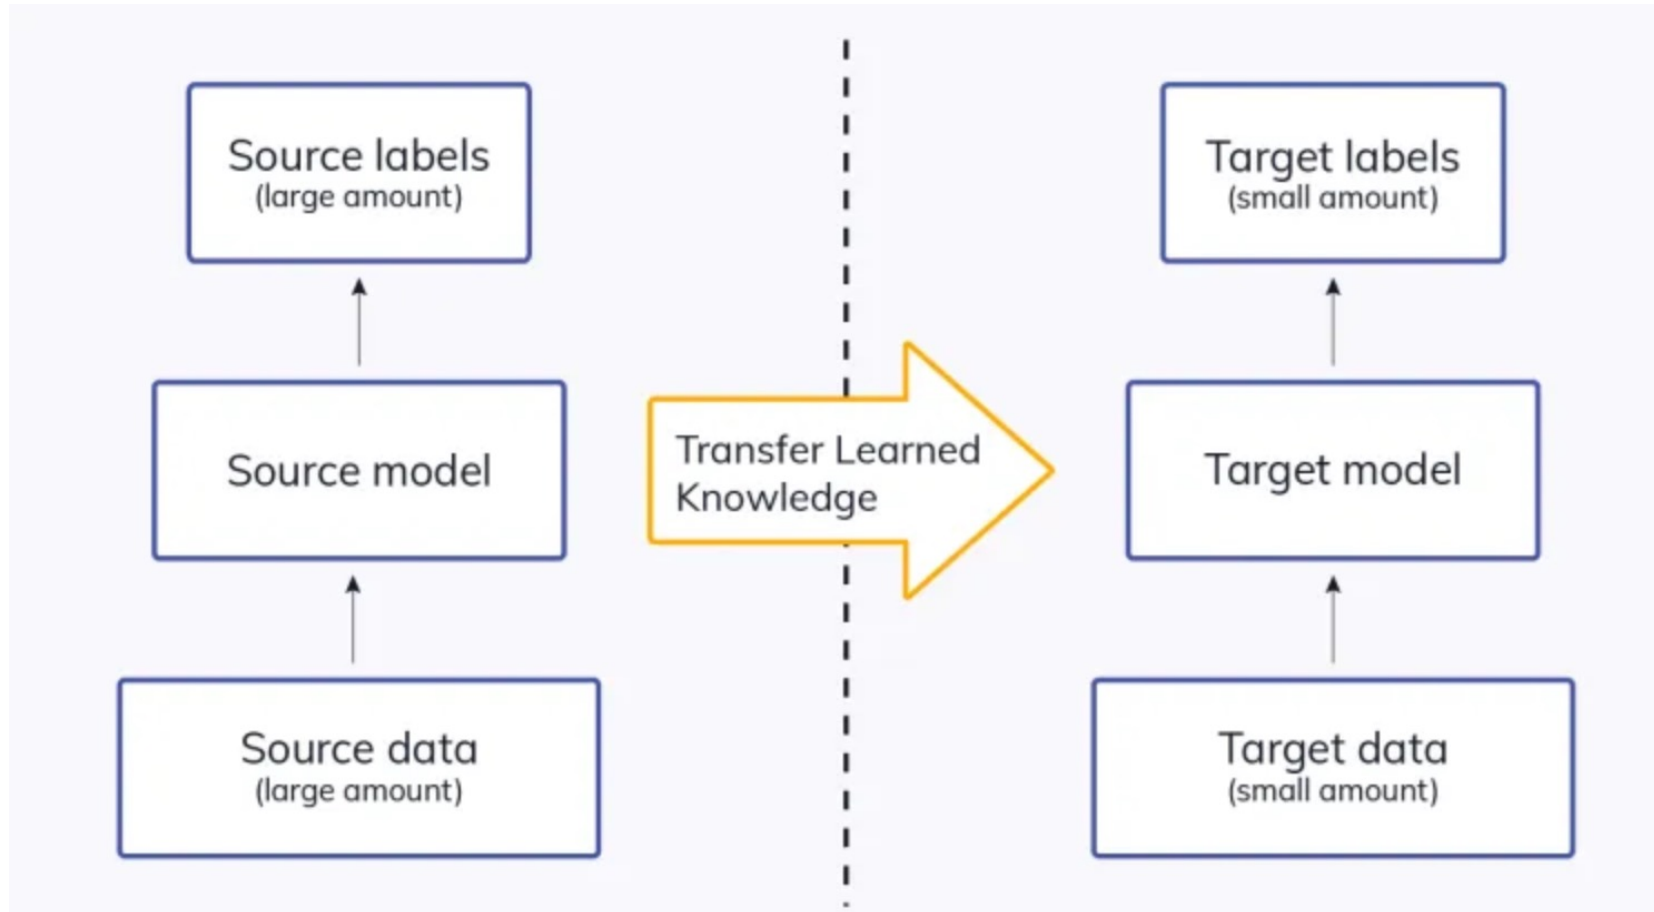
- Contextual understanding
- Bi-directional learning

- **Cons:**

- Computational complexity
- Lack of interpretability
- Large memory footprint

Transfer Learning and NLP Applications

Transfer Learning : Idea




Transfer Learning


- Supervised learning for one task and transferred to another task that does not have a lot of labeled data
- Typically, models are pre-trained on large amount of data for a well known task
- Transferred to other task using small amount of training data for the target task
- Highly popular in Deep Learning

Transfer Learning in text - Example

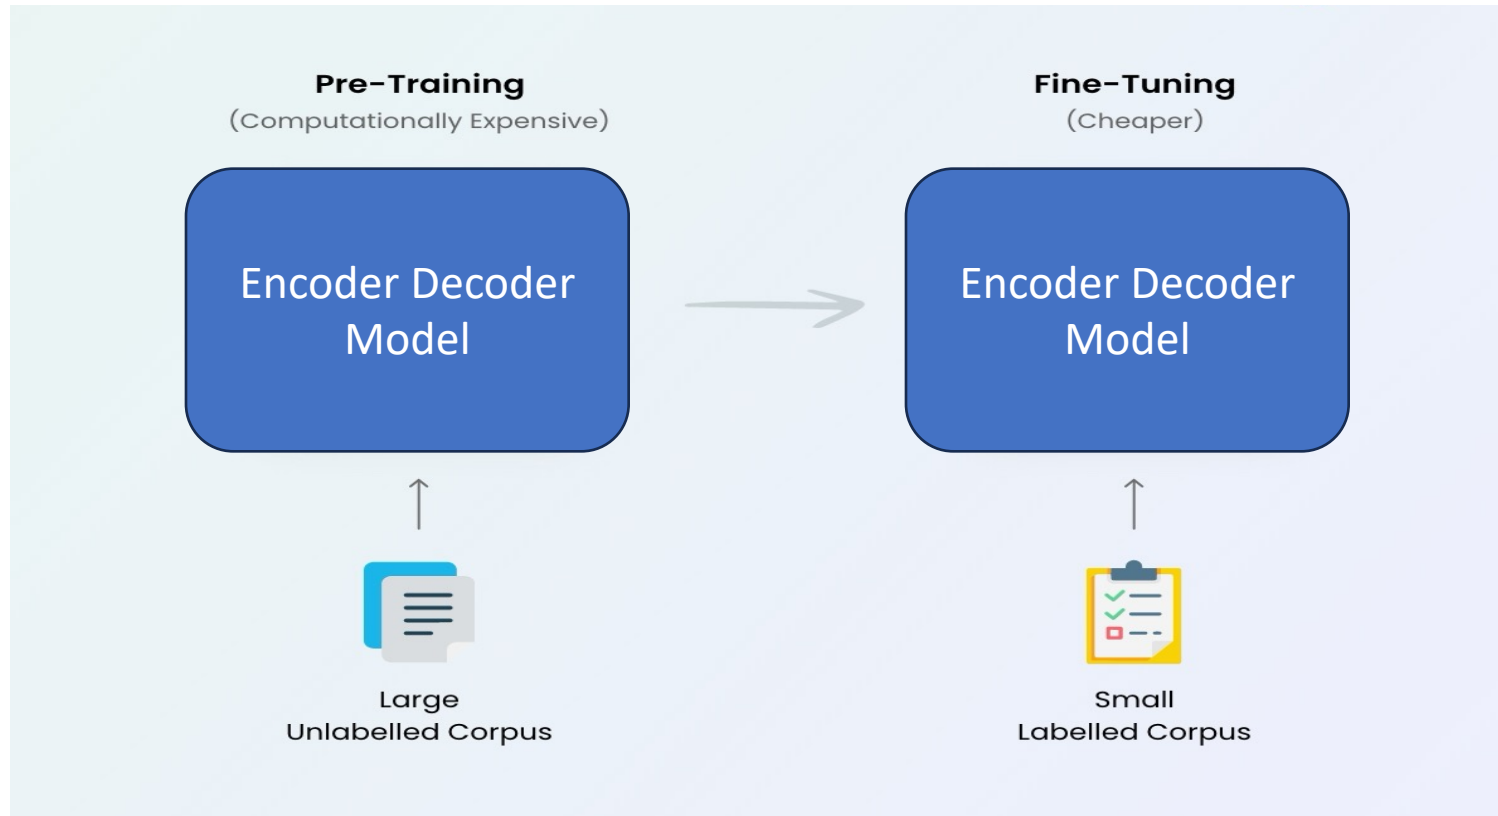
- Train a language model for reconstruct Missing. / Future text.



- Transfer and fine-tune “a branch” of the trained model to do specialized text generation (e.g., Poem), with few examples

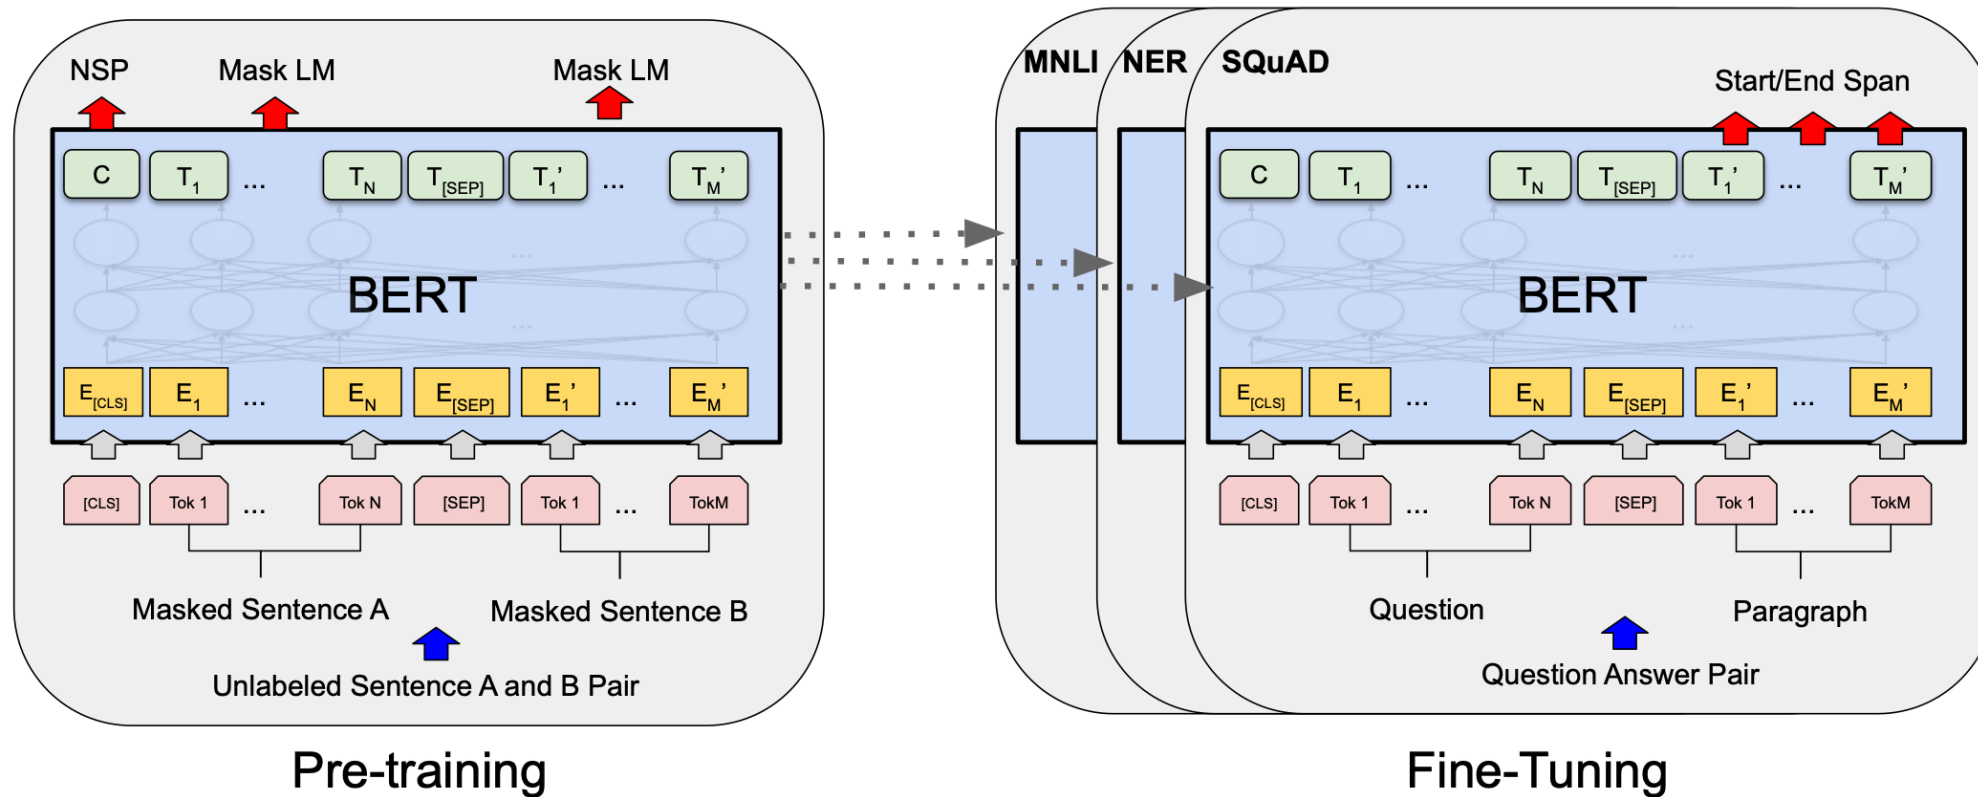


Transfer Learning in Text



Some Pre-trained Language Models

BERT



BERT – pretraining

- BERT aims to learn a contextualized representation for each token in a given sentence by optimizing the following objective function
- Let T be the input sequence of tokens. BERT learns the parameters θ by maximizing the log likelihood of the next sentence prediction and the masked language model objectives

$$\max_{\theta} \sum_{(T_a, T_b) \in D} \log p(IsNext|T_a, T_b, \theta) + \sum_{t \in T} \log p(t|T_{-t}, \theta)$$

Dataset used

- BERT was trained on a large corpus that includes BooksCorpus (800 million words) and English Wikipedia (2,500 million words).

Tokenization

- BERT uses WordPiece tokenization, which breaks words into subwords based on a fixed vocabulary. This approach enables the model to handle out-of-vocabulary words effectively.



Downstream Fine-tuning Tasks:

- BERT can be fine-tuned for various downstream NLP tasks, including but not limited to text classification, question-answering, named entity recognition, text entailment, and sentiment analysis.

RoBERTa

- An extension of BERT

Optimizes only Masked LM objective

$$\max_{\theta} \sum_{t \in T} \log p(t | T_{\neg t}, \theta)$$

Tokenization

- Similar to BERT, RoBERTa utilizes WordPiece tokenization for subword token handling.
-

Dataset


- RoBERTa was trained on a combination of in-domain data (books and articles) and out-of-domain data (web data).

GPT


- The objective is to maximize the log likelihood of the next token in the sequence:

$$\max_{\theta} \sum_{t \in T} \log p(t | T_{<t}, \theta)$$

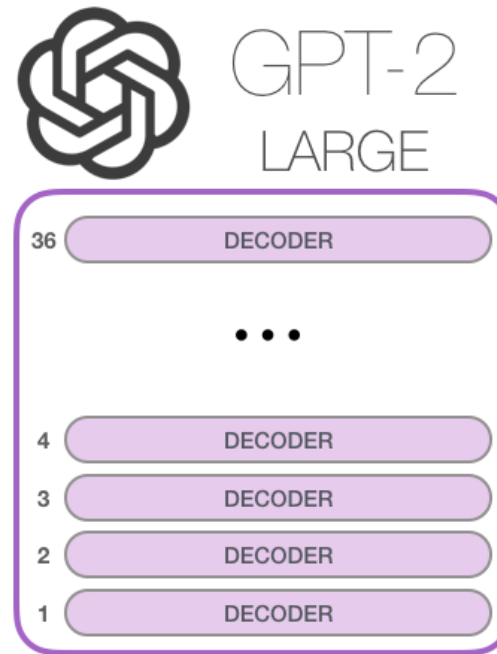
GPT Evolution



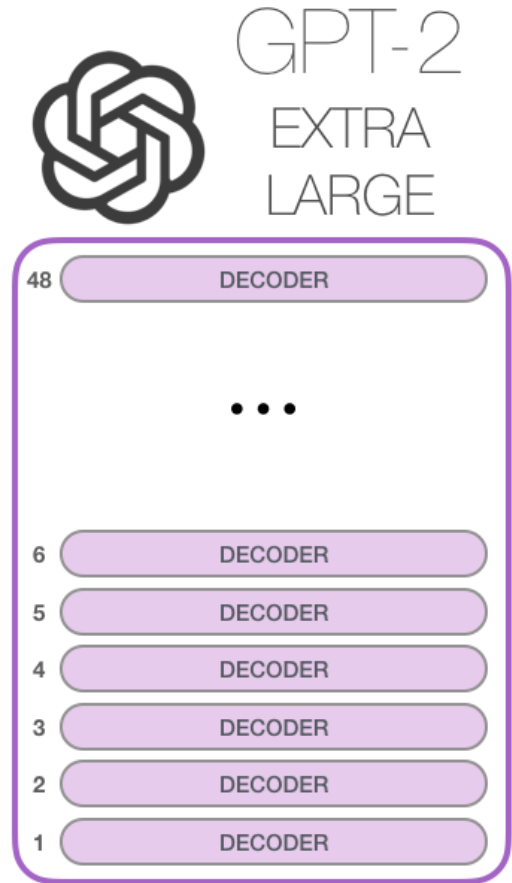
Model Dimensionality: 768



Model Dimensionality: 1024



Model Dimensionality: 1280



Model Dimensionality: 1600

Tokenization and Dataset

- GPT-2 uses byte pair encoding (BPE) for handling subword tokenization, allowing the model to handle rare and unseen words efficiently.
- GPT-2 was trained on a diverse range of internet text data, encompassing a wide array of sources to ensure a broad understanding of human language.
- Other higher order GPTs follow similar tokenization . Datasets are not disclosed.

BART

- BART is trained as a denoising autoencoder, where the model is tasked with reconstructing the original text from a corrupted version. Its objective is to minimize the reconstruction error, which can be formulated as:

$$\min_{\theta} \sum_{T \in D} \mathbb{E}_{\tilde{T} \sim \text{Corrupt}(T)} [-\log p(T | \tilde{T}, \theta)]$$

~

Tokenization and Dataset

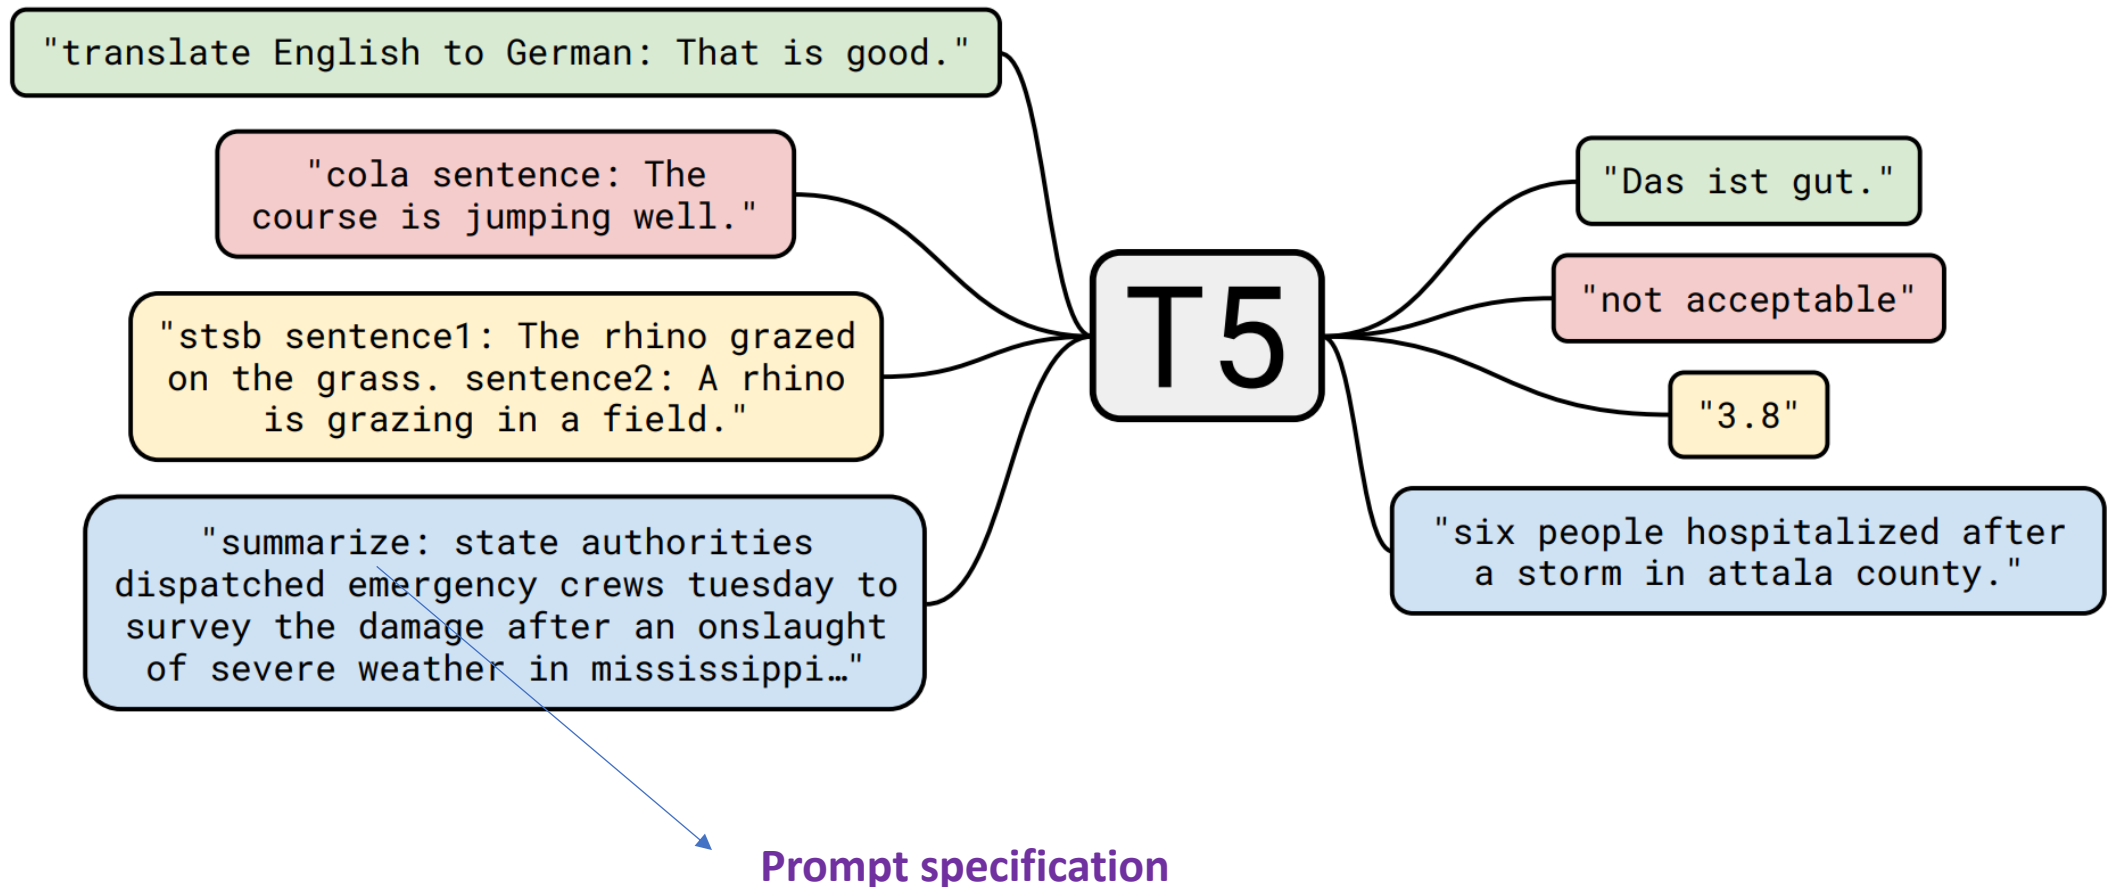
- BART employs a combination of Byte Pair Encoding (BPE) and learned positional embeddings to handle tokenization.
- BART was trained on a mixture of text from a wide range of sources, including books, articles, and websites.

T5

- T5 formulates all tasks as text-to-text problems, unifying different NLP tasks into a single framework. The objective function is to maximize the log likelihood of the target text given the input text:

$$\max_{\theta} \sum_{(X,Y) \in D} \log p(Y|X, \theta)$$

T5: Treating all problems as a language modeling task



Tokenizers and datasets

- T5 employs a variant of the Byte Pair Encoding (BPE) algorithm for subword tokenization.
- T5 was trained on a diverse corpus, including books, articles, and websites, ensuring a broad understanding of human language.

Fine tuning pre-trained models

- Fine-tuning involves taking a pre-trained model and adapting it to a specific task
- Generic steps:
 - **Select Pre-Trained Model:**
 - E.g. `bert-base-uncased` for lower-cased data, `bert-base-cased` for true cased data
 - `bert-base-multilingual` for multilingual tasks (e.g., Translation)
 - **Data Preparation**
 - **Model Architecture Modification (if necessary)**
 - **Initialize Parameters with the pretrained model**

Fine tuning pre-trained models

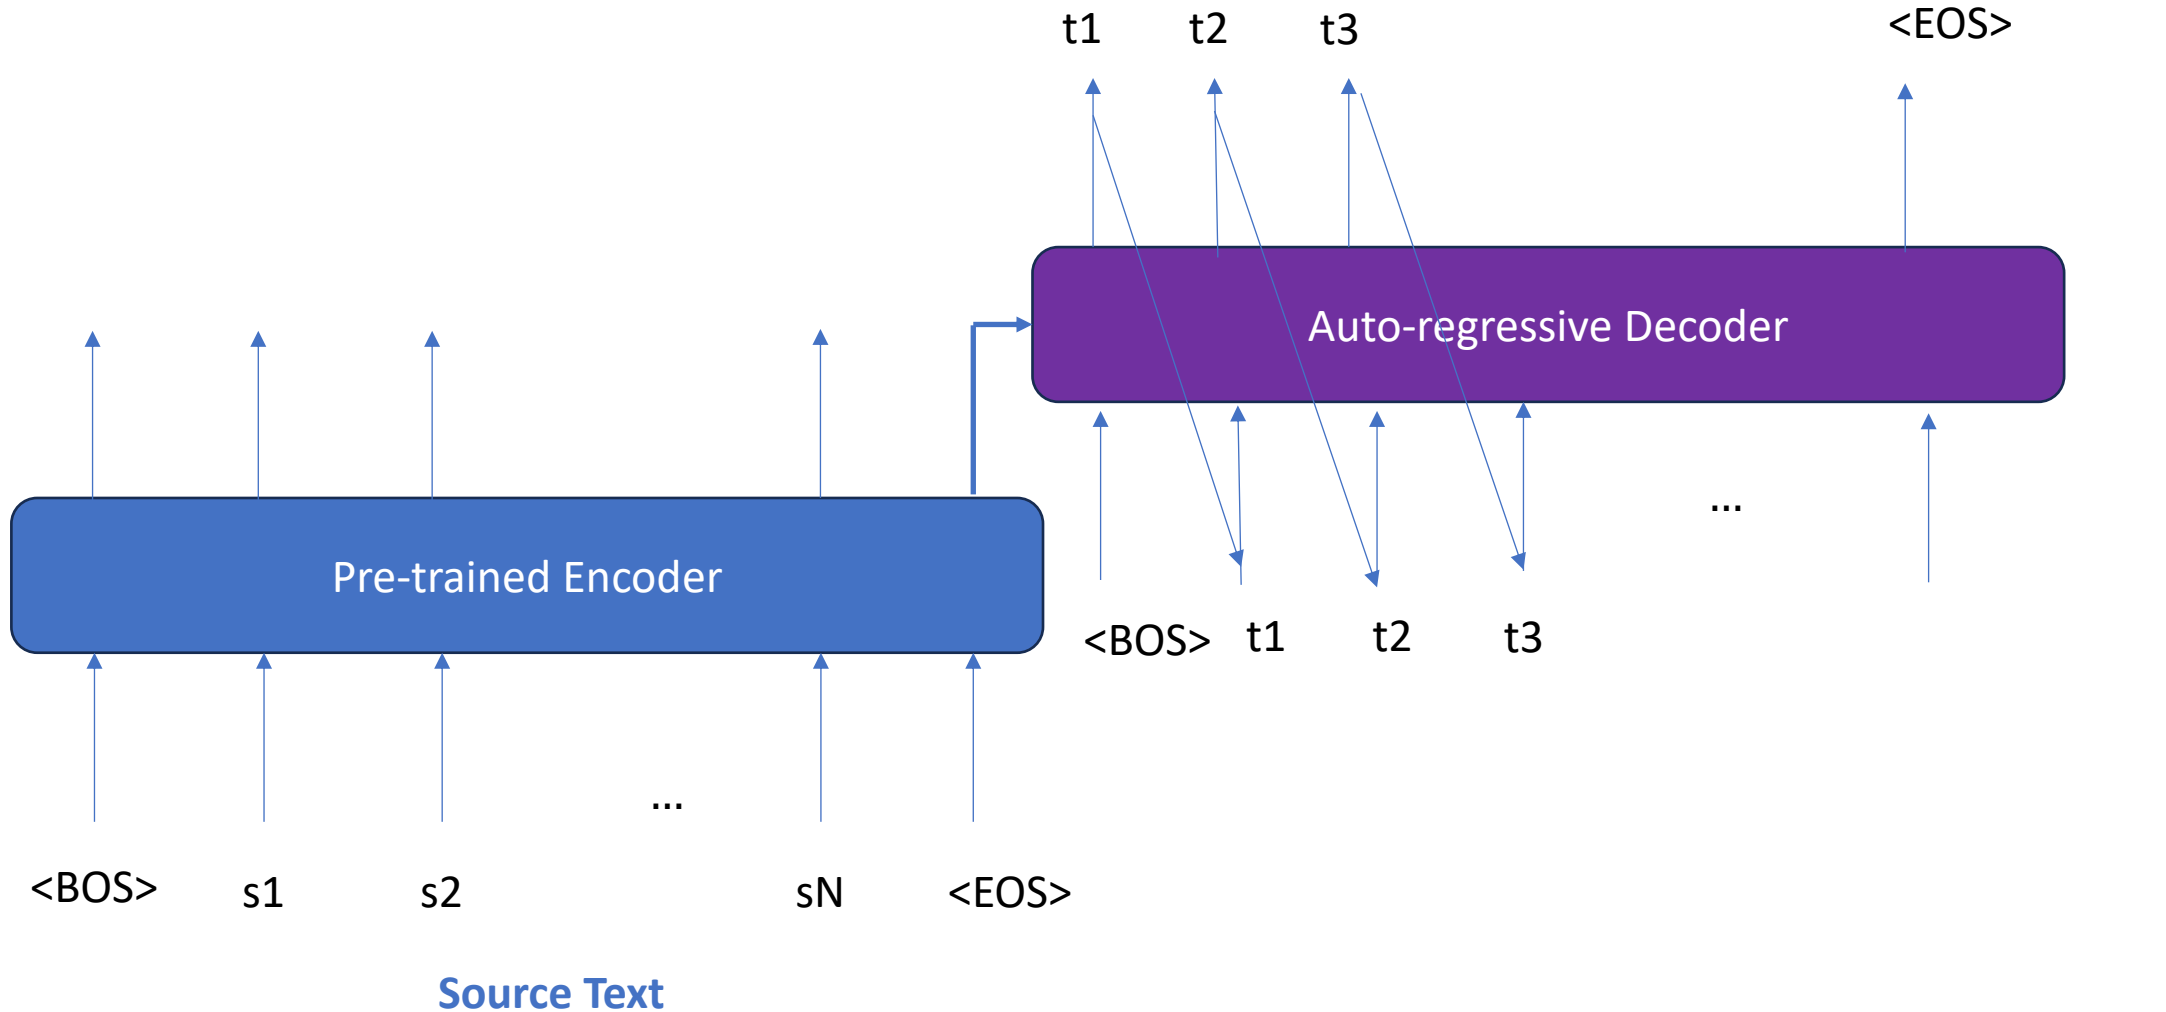
- Fine-tuning involves taking a pre-trained model and adapting it to a specific task
- Generic steps:
 - **Fine tuning Process:** Train the model on the task-specific dataset while monitoring its performance. Print train and validation loss
 - **Hyperparameter Tuning (if necessary):** Fine-tune the hyperparameters if the model's performance is not satisfactory. This might involve adjusting the learning rate, batch size, or other optimization parameters.
 - **Testing and Deployment**

Downstream Tasks – Machine Translation

- **Machine Translation**

- The T5 model by Google has been widely used for various NLP tasks, including machine translation. It can be fine-tuned for specific translation tasks, making it an effective choice for this purpose
- **Fairseq** provides pre-trained models like **Transformer**, **Transformer Big**, and **Transformer WMT19**, which are commonly used for machine translation tasks.

Typical architecture – seq2seq



Downstream Tasks – Machine Translation

- **Datasets:**

- WMT (Workshop on Machine Translation) datasets, including WMT14, WMT16, and WMT19.
- IWSLT (International Workshop on Spoken Language Translation) datasets.
- Multi30k dataset.
- TED Talks dataset.

- **Evaluation Metrics:**

- BLEU (Bilingual Evaluation Understudy): Measures the quality of machine-translated text by comparing it to one or more reference translations.
- METEOR (Metric for Evaluation of Translation with Explicit Ordering): Considers unigram matching, stem matching, and synonymy.

Downstream Tasks – Summarization

- **Datasets:**

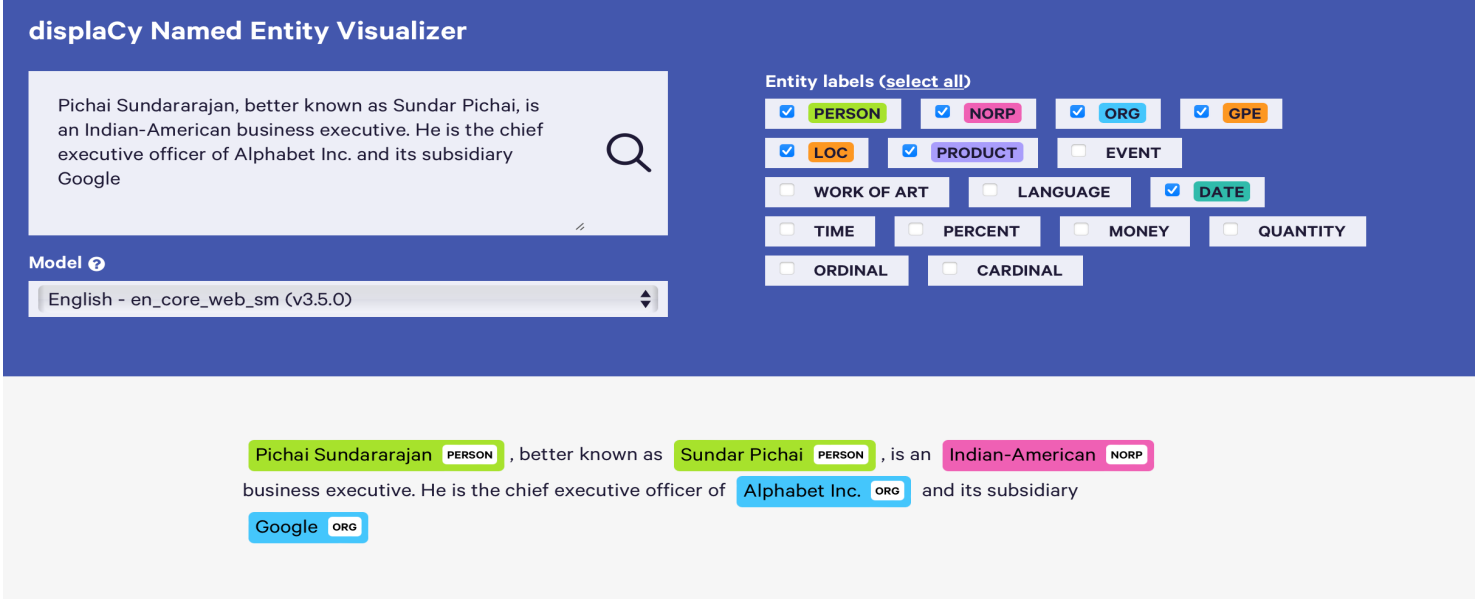
- CNN/Daily Mail dataset.
- XSum dataset.
- Gigaword dataset.
- Newsroom dataset.

- **Evaluation Metrics:**

- **ROUGE (Recall-Oriented Understudy for Gisting Evaluation):** Measures the overlap between the generated summary and the reference summaries at different levels (unigram, bigram, etc.).
- **BLEU (Bilingual Evaluation Understudy):** Often used to evaluate the quality of generated summaries by comparing them to one or more reference summaries.

Sequence Tagging –

- **Tagging each input token with a class label**
 - **Example:** Part of Speech tagging
 - Named Entity Recognition



The screenshot displays the displaCy Named Entity Visualizer interface. On the left, a text input field contains the sentence: "Pichai Sundararajan, better known as Sundar Pichai, is an Indian-American business executive. He is the chief executive officer of Alphabet Inc. and its subsidiary Google". Below the input is a search icon and a model selector dropdown menu set to "English - en_core_web_sm (v3.5.0)". On the right, a panel titled "Entity labels (select all)" lists various entity classes with checkboxes. The selected labels are PERSON, NORP, ORG, GPE, LOC, and PRODUCT. Below the list, the visualized output shows the original sentence with colored boxes and labels identifying entities: "Pichai Sundararajan" (PERSON), "Sundar Pichai" (PERSON), "Indian-American" (NORP), "Alphabet Inc." (ORG), and "Google" (ORG).

displaCy Named Entity Visualizer

Pichai Sundararajan, better known as Sundar Pichai, is an Indian-American business executive. He is the chief executive officer of Alphabet Inc. and its subsidiary Google

Model ?

English - en_core_web_sm (v3.5.0)

Entity labels (select all)

☒ PERSON ☒ NORP ☒ ORG ☒ GPE

☒ LOC ☒ PRODUCT ☐ EVENT

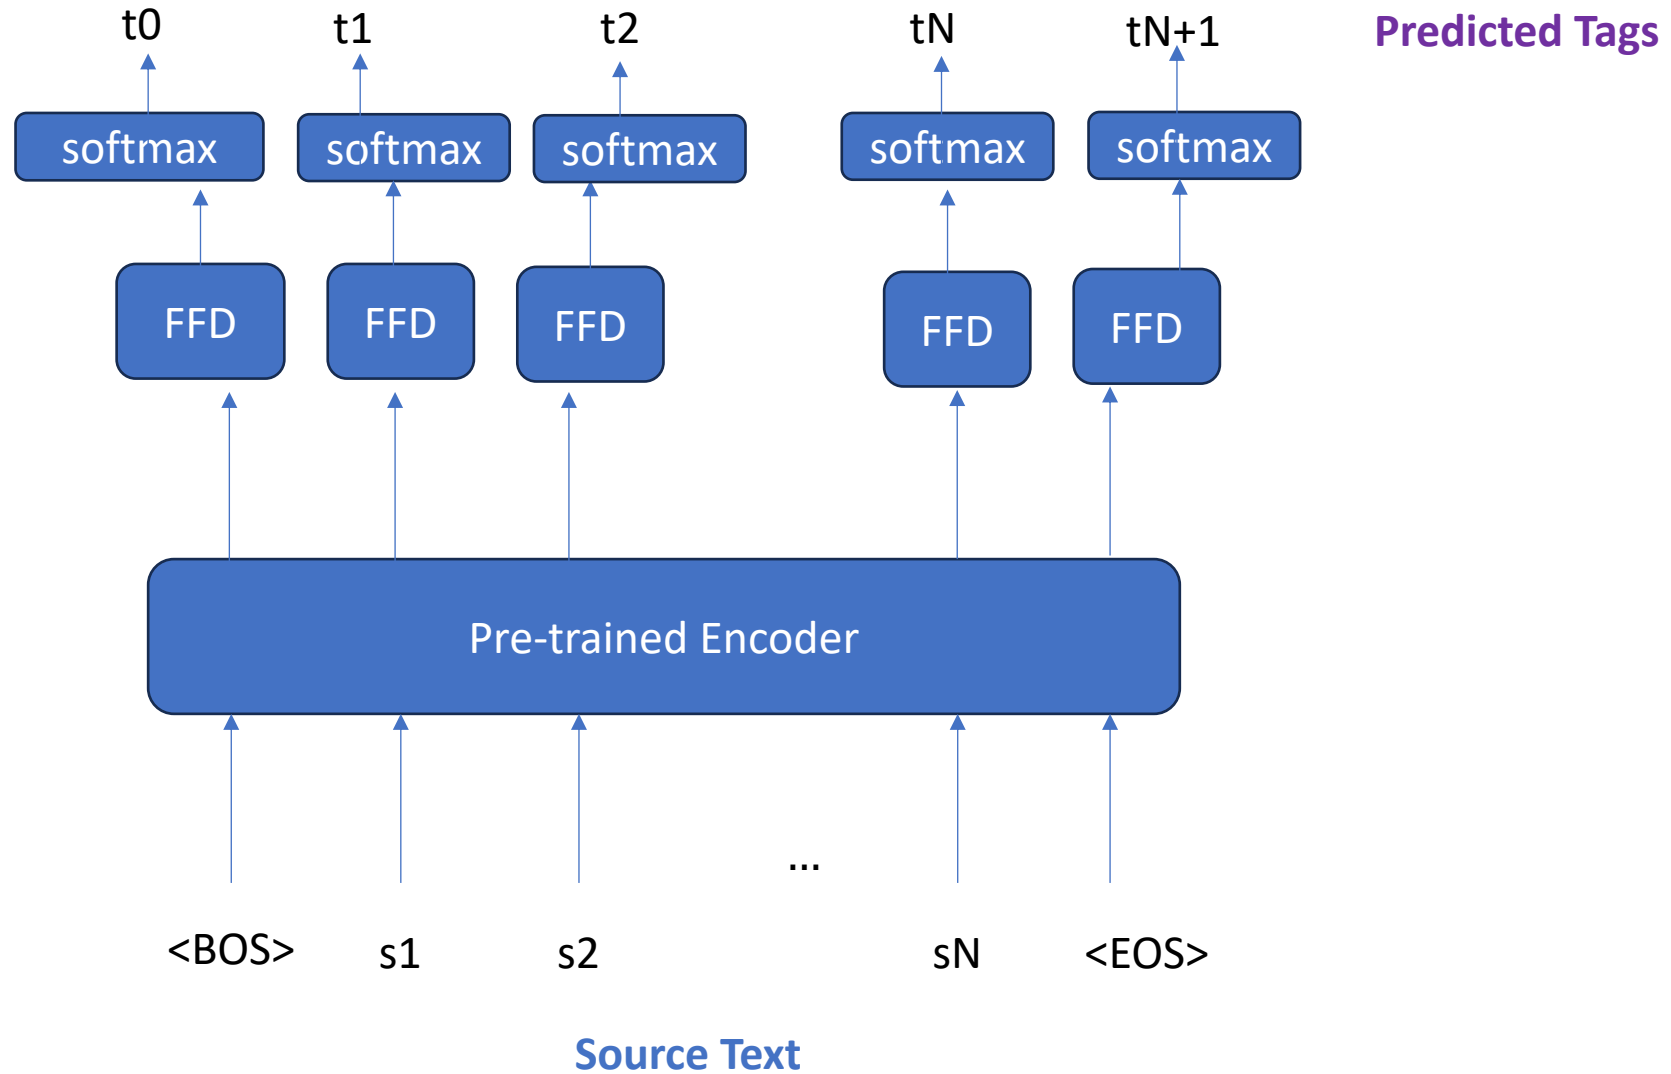
☐ WORK OF ART ☐ LANGUAGE ☒ DATE

☐ TIME ☐ PERCENT ☐ MONEY ☐ QUANTITY

☐ ORDINAL ☐ CARDINAL

Pichai Sundararajan PERSON, better known as Sundar Pichai PERSON, is an Indian-American NORP business executive. He is the chief executive officer of Alphabet Inc. ORG and its subsidiary Google ORG

Typical architecture – sequence labeling



Sequence Tagging – NER

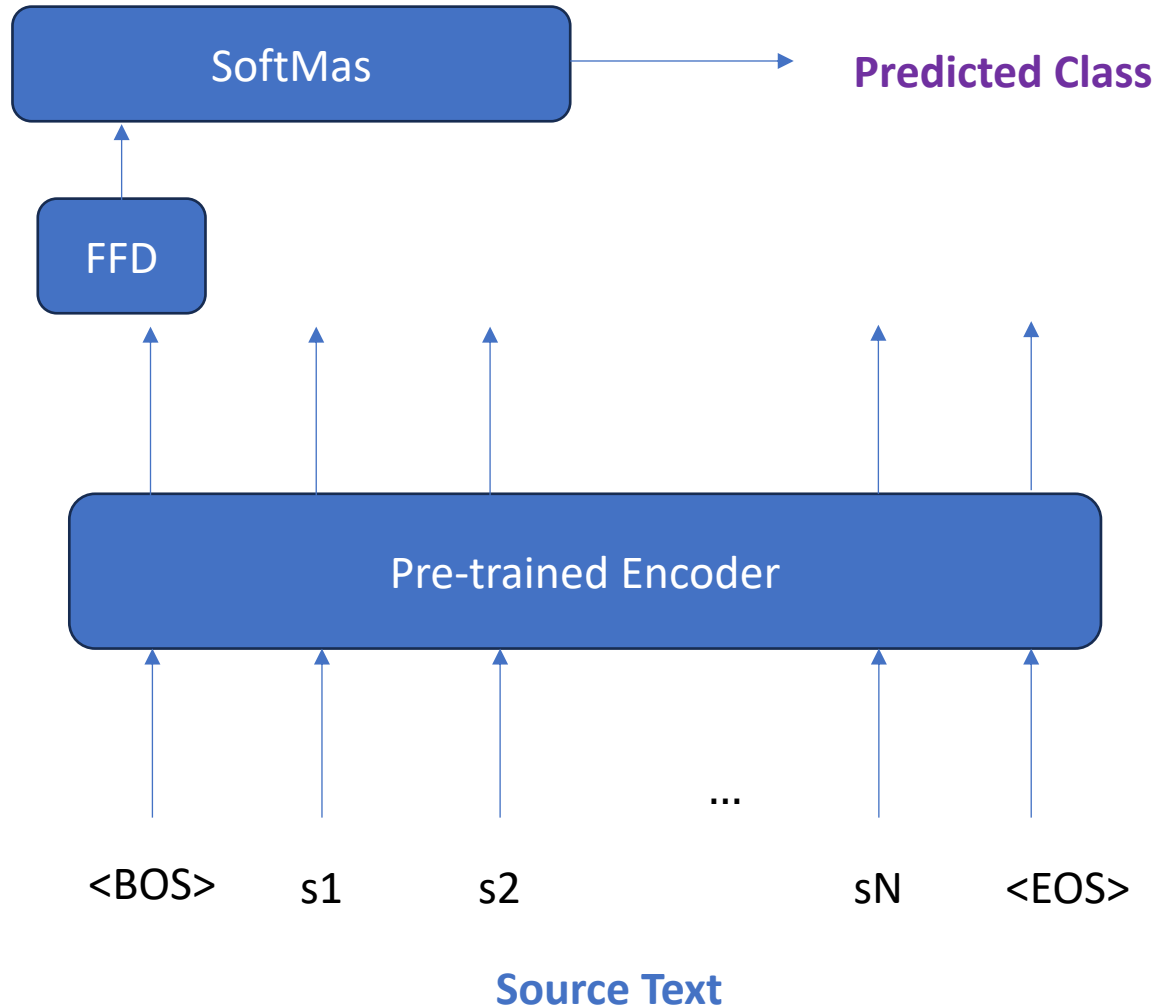
- **Datasets:**

- CoNLL 2003 dataset.
- OntoNotes dataset.
- GermEval dataset.
- ACE (Automatic Content Extraction) dataset.

- **Evaluation Metrics:**

- Precision, Recall, and F1-score: Commonly used to evaluate the performance of named entity recognition systems by comparing the predicted entities to the ground truth entities.
- CoNLL score: Used to evaluate the overall performance of a named entity recognition system, combining precision and recall into a single metric.

Typical architecture – text classification



Considering only the context vector from first input is enough

Text Classification

- **Datasets:**

- IMDB Movie Reviews dataset.
- AG News dataset.
- Yelp Reviews dataset.
- DBpedia dataset.

- **Evaluation Metrics:**

- Accuracy: Measures the proportion of correctly classified instances.
- Precision, Recall, and F1-score: Used to evaluate the performance of the classification model, particularly in tasks where class imbalance is present.

Next class

Fine-tuning Language Models for Building Classification and Generation Systems